

© 2025 Frank Andabwa

All rights reserved.

No part of this publication may be reproduced, stored in a retrieval system, or transmitted in any form or by any means, electronic, mechanical, photocopying, recording, or otherwise, without the prior written permission of the author, except in the case of brief quotations embodied in critical articles or reviews.

TABLE OF CONTENTS

- **Executive Summary**
- **Introduction**
- 1. Foundations of Macro ETF Strategy Design
 - 1.1 A Top-Down, Macro-Quant Approach
- 2. Designing the Hedge Engine Strategy
 - 2.1 Identifying Macroeconomic Signals
 - 2.2 Sector Rotation Across Business Cycles
 - 2.3 Event-Driven Triggers and Probabilistic Timing
 - 2.4 Volatility Targeting and Synthetic Leverage
 - 2.5 Synthesis of Strategy Rules (Six-Step Execution Flow)
- 3. Backtesting and Performance Evaluation
 - 3.1 Backtest Framework:
 - 3.2 Interpreting Results:
- 4. LLM Integration: Crafting the Strategy as a Meta-Prompt
 - 4.1 Clear Structure and Tags
 - 4.2 Meta-Prompt Sidebars
 - 4.3 Prompt Blueprint Example
 - 4.4 Format for LLM Parsing
 - 4.5 Continual Learning and Adaptation
- 5. Risk Management and Disclosure
 - 5.1 Key Risk Factors:
 - 5.2 Disclosures (for a Professional Document)
 - 5.3 Risk Management Practices
- 6. Implementation and Rollout Plan
 - 6.1 Phase 1: Preparation and Infrastructure
 - 6.2 Phase 2: Paper Trading and Calibration
 - 6.3 Phase 3: Live Rollout
 - 6.4 Data Sources and Tools - Summary List
 - 6.5 Ongoing Maintenance and Evolution
- **Conclusion**

Addendums (Definitive Policy and Implementation)

- **Addendum I: Investment Parameters and Limits**
 - Target Annual Volatility
 - Volatility Measurement Window
 - Max Scaling Factor (Leverage)
 - Minimum Exposure Floor
 - Base Rebalancing Frequency
 - Emergency Drawdown Trigger
 - Max Single Sector Weight
- **Addendum II: Definitive ETF Universe and Tickers**
 - Broad Equity Market
 - Cyclical Sectors (Tech, Discretionary, Industrials, Financials)
 - Defensive Sectors (Staples, Healthcare, Utilities)
 - Long-Term Bonds
 - Inflation/Commodity Hedge
 - Cash Proxy / Short-Term Bonds
 - Approved Leveraged Instruments
- **Addendum III: Backtest Performance Summary (Placeholder)**
- **Addendum IV: Auditable Functionality**
 - Principles (non-negotiable)
 - Required Components (Decision Record JSON Schema, etc.)
- **Addendum V: Custom GPT Configuration**
- **Addendum VI: Operator Instructions: Make Hedge Engine Agentic**

Hedge Engine: A Macro ETF Strategy Guide and LLM Blueprint

Executive Summary

I built Hedge Engine because I could no longer stand watching the same story repeat itself: hardworking people thinking they were invested sensibly, then getting blindsided by market shocks that ruin retirements and erase a lifetime of savings. That realization was the blunt force that created Hedgehog, and it is the moral engine behind everything I do now.

The problem I set out to solve is simple and urgent. The average retail investor is underprotected against extreme market events. History proves this again and again. On Black Monday in 1987 the market dropped roughly twenty two percent in a single day. In 2010 the Flash Crash erased nearly ten percent of market value inside minutes. The 2008 financial crisis was far worse than a correction; it showed that products marketed as de risked can still inflict catastrophic losses on ordinary people. In 2008 many funds that were advertised as conservative or de risked lost money for the vast majority of participants, and average losses were, in many cases, devastating.

Those events exposed a structural flaw. Conventional portfolio theory and the models used by many retail platforms and robo advisors assume a smooth, bell shaped world where risk behaves politely. The real market is not polite. Over the past three decades market behavior has repeatedly invalidated those assumptions. Markets move in sharp, path dependent ways, with regime shifts and nonlinear dynamics. Treating risk as a tidy number that fits into a spreadsheet is not a minor modeling error. It is a design failure that harms people.

At the same time, too many households are exposed. Roughly half of U.S. households have some stock market exposure, either directly or through retirement accounts. A large share of those investors do not understand the full scale of the tail risk they carry, and most do not have a coherent, operational plan to mitigate it. That exposure gap is the problem we had to solve: people are invested, but they are not protected.

I came to this from FinTech and big data, and the gap felt visceral. I did not want another abstract paper that told investors what might happen without giving them a practical tool. I wanted protection that was operational, repeatable, and verifiable. That was the conceptual leap that produced Hedgehog: an investor risk engine designed to prevent terrible outcomes for ordinary investors while improving risk adjusted returns where possible.

Turning that idea into reality meant assembling a stack of technologies and a clear product model. I combined four main elements.

1. Big data to profile real investor risk. The usual approach treats risk as a questionnaire answer, a single number. I built a richer risk profile. We map liquidity needs, time horizons, behavioral tendencies, portfolio concentration, and scenarios that would be ruinous for each household. That profile is the input that determines what the engine protects against.
2. Artificial intelligence to model markets honestly. I train and deploy models that explicitly represent path dependence, regime shifts, and non linear propagation of shocks. These models do not pretend markets follow Gaussian curves. They are designed to generate probabilistic scenarios that feed surgical hedges and tactical tilts.
3. Crowdsourced human intelligence as a safety valve. Machines scale and detect patterns. Humans still spot novel edge cases, geopolitical nuance, and legal or policy subtleties that raw statistics miss. We structured a human in the loop to validate unusual signals and to audit model outputs where context matters.
4. A blockchain layer where integrity and verification matter. I use blockchain selectively as an auditable, tamper resistant ledger for model outputs, evidence citations, and governance events. This is not a gimmick. When you automate risk decisions at scale, you must be able to prove what the system

saw and why it acted.

The product idea was also important. I did not want to build a closed hedge fund or an exclusive tool for institutions. I envisioned Robo Advisor 2.0, a platform that behaves like an easy button for risk. Think of it as a search bar for risk: you enter a user context, and the platform returns contextual, market fitting recommendations that integrate into the platforms people already use. The goal is practical automation and broad access, not mystique.

Hedge Engine is the operationalization of that vision. I translated an abstract risk engine into a codified blueprint made of three institutional grade overlays. Each overlay is crisp, rule based, auditable, and designed to work together.

First, macro driven sector rotation. Passive, static allocation is convenient but dangerous in the wrong regime. If you hold the same mix while the economy rotates, you can be the last investor left standing in the wrong sector when the cycle turns. So we infer the business cycle from a basket of indicators, such as GDP nowcasts, purchasing managers indexes, the yield curve, credit spreads, and employment trends. Then we tilt sector exposure according to the regime. Early expansion calls for cyclical overweight, mid cycle favors broad diversification, late cycle calls for defensive tilts into sectors like health care, consumer staples, and utilities, and recession calls for capital preservation with Treasury exposure and defensive assets. The rotation is rule based and gradual to avoid whipsaw. It

is a framework for being in the right place when the cycle turns and for avoiding being caught flat footed.

Second, event driven triggers. Markets move not only by slow cycles but via sharp events, both scheduled and surprise. For scheduled events, such as central bank meetings, major data releases, or elections, Hedge Engine models outcomes, assigns probabilities, and sizes hedges proportionally to both probability and impact. That is probabilistic timing: we do not overreact to every headline, and we do not pretend to forecast with certainty. We scale exposure, where appropriate, according to the odds and the expected market response. For surprise events, we deploy market sensors. Spikes in volatility, sudden credit spread widening, or liquidity evaporation trigger circuit breakers that rapidly reduce exposure. The objective is systematic, graded responses, not emotional overreaction.

Third, volatility targeting. This is the ultimate defense layer. We do not pretend to eliminate volatility; we manage exposure to it. The engine measures realized volatility and dynamically adjusts the portfolio to maintain a target level of risk. When markets are calm, we can scale up exposure to reach the target. When markets become turbulent, we scale down to preserve capital. That produces a steadier risk budget over time, fewer catastrophic drawdowns, and a smoother investor experience. Implementation includes explicit scaling rules, practical limits on leverage, and guardrails to prevent over leveraging in calm regimes. Where we use synthetic leverage and hedging, we do so with operational discipline and explicit rules to manage the decay and margin call risks that naive leverage produces.

I designed these overlays to be precise. Each one is codified as rules and edge cases rather than vague recommendations. That was intentional. The final blueprint is written to be both human readable and machine executable. It functions as a meta prompt for a large language model: the LLM receives the decision logic, the evidence requirements, the thresholds, the fail safes, and the audit trail expectations. The goal is operational discipline. An LLM executing the blueprint will not panic in a crisis, it will not guess without evidence, and it will create an auditable record of why it acted. That combination of machine scale, computational precision, and human governance is the practical hedge that retail investors have lacked.

A few additional practical details about how this works in operation. For macro driven rotation, we smooth signals and avoid flipping positions overnight. We favor liquid sector ETFs to express tilts and use cash or short term Treasuries as risk sinks during stormy periods. For event driven triggers, we combine market implied probabilities, option market skew, and fundamental signals to calibrate hedges. For volatility targeting, we use rolling volatility estimates, a target volatility band, and scaling factors that are capped so we never exceed prudent leverage limits. Throughout, we log the evidence behind each decision: which indicators moved, which documents or data points the model used, and who, if anyone, validated the action.

Finally, I built governance into the product from day one. Every decision path is auditable. Every model update and prompt change is versioned. We maintain human oversight

and clear escalation rules. We treat the hedge as a public good: the system is designed to protect households that rely on it, and to be explainable to regulators, auditors, and the people who actually put money into these strategies.

In short, Hedge Engine is not a marketing tagline. It is the operational answer to a real social problem: households exposed to catastrophic market risk without the tools to manage it. Hedge Engine profiles true risk, defends against the things models miss, and gives a practical, repeatable playbook you can implement and audit. I built it to stop terrible things from happening to good investors, and I built it so the protections are real, documented, and repeatable in the messy, chaotic market environments where people need them most.

Introduction

Key Points: This guide presents a **macro-driven ETF investment strategy** that adapts to economic cycles and events, and it is structured as a “**meta-prompt**” meaning the content is formatted for clarity to humans *and* as a functional blueprint for a Large Language Model (LLM) to execute the methodology. We will cover the design of a **hedge-fund-style strategy using ETFs**, including macroeconomic signals (e.g. business cycle indicators), event-driven triggers (like central bank decisions), and volatility/leverage overlays. Throughout, we employ professional formatting conventions from institutional whitepapers, clear section headings, concise bullet steps, and data-driven charts to make the material both accessible

and algorithmically parseable. **Important:** This document is for educational purposes; it is *not* personalized investment advice. All investment strategies carry risk, especially those involving leverage or complex instruments, and past performance is not indicative of future results.

This *Hedge Engine* ebook is organized into logical chapters that build the strategy step by step. Human readers will find explanatory sections with real-world examples, while sidebars and pseudo-code snippets translate those concepts into explicit rules that an LLM or algorithm could follow. By the end, you should understand how to construct a macro-ETF “hedge fund” strategy and have a template for how an AI could implement and manage it in real time. Let’s begin by laying the foundation of macroeconomic strategy design.

1. Foundations of Macro ETF Strategy Design

In contrast to security-specific investing, **macro strategies** focus on broad economic trends and events to drive investment decisions. Exchange-Traded Funds (ETFs) serve as flexible building blocks for these strategies, giving even a retail investor targeted exposure to asset classes, sectors, or themes. A well-designed macro ETF strategy can resemble a hedge fund in its **diversification** and dynamic allocation, but with transparency and lower costs. The *Hedge Engine* approach will systematically rotate between ETFs (equities, bonds, commodities, sectors, etc.) based on macro signals and risk conditions, aiming to

hedge downside risks and capture upside across different market regimes.

To ensure clarity and credibility, we model the formatting here on institutional investment reports. That means using an **organized structure** (e.g. numbered sections and sub-sections), **short paragraphs** for readability, and visual aids like charts and tables for key insights. Professional white papers often begin with an abstract or key points (as we have) and use formal yet accessible language. We adopt that style so that complex ideas (like volatility targeting or probabilistic event timing) are presented in digestible form. This structured approach not only benefits human readers but also helps an LLM break down the methodology into components. Each major concept will be clearly labeled and often accompanied by bullet-point steps or pseudo-code illustrating how to implement it.

1.1 A Top-Down, Macro-Quant Approach

Hedge Engine's strategy is **top-down**: it starts with the big picture: economic growth, inflation, interest rates, policy changes, and then determines investments accordingly. This is akin to what macro hedge funds do, but implemented through liquid ETFs. Some foundational principles:

- **Macro Indicators Drive Decisions:** Rather than picking stocks on valuation or micro factors, we use **macroeconomic indicators** (e.g. GDP growth rates, PMI indices, yield curve slope, unemployment trends) to infer which assets are likely to perform

well. For example, if leading indicators suggest a recession is coming, the strategy might tilt toward Treasury bond ETFs or defensive equity sectors.

- **ETF Selection for Targeted Exposure:** Each macro view is expressed via ETFs. Want to bet on rising inflation? Perhaps use a commodity or inflation-protected bond ETF. Expecting an economic expansion? Tilt toward cyclical sector ETFs (technology, consumer discretionary, industrials, etc.). ETFs allow quick rotation in and out of themes without needing to pick individual winners.
- **Risk Management at the Core:** A macro strategy must manage risk actively, our approach will include **volatility targeting** (adjusting exposure to maintain steady risk) and possibly using **leveraged ETFs or derivatives** in a controlled way (what we term *synthetic leverage*). This can amplify returns but also magnifies losses, so clear rules are vital. We will incorporate common risk controls used by professional funds, like drawdown limits and diversification rules.
- **Event Responsiveness:** The strategy isn't purely systematic on economic data; it also has an **event-driven** element. Major scheduled events (Federal Reserve meetings, jobs reports, elections) and even unexpected shocks (financial crises, geopolitical flare-ups) can drastically shift market dynamics. Hedge Engine will define trigger

conditions for certain events for instance, reducing equity exposure before an anticipated rate hike, or reallocating to safe havens if a geopolitical risk escalates. We use *probabilistic timing*, meaning we assess the probability of events (e.g. using Fed Funds futures to gauge odds of a Fed hike) and adjust positions in proportion to those probabilities.

- **Structured for Humans and Machines:** Throughout this guide, you'll see **Meta-Prompt Sidebars** and **Pseudo-Code** segments. These are special callouts where we translate the discussed strategy logic into structured rules or code-like language. For a human reader, they serve as a summary or clarification. For an LLM, these sidebars act as a blueprint for turning the described methodology into actual operations. In a sense, the entire ebook is a meta-prompt, a set of instructions and context that an AI can use to simulate the hedge fund manager's decision-making.

Before diving into the mechanics, let's outline the design of the strategy, identifying which macro signals and frameworks we'll use.

2. Designing the Hedge Engine Strategy

In this chapter, we build the strategy's components: the economic signals to track, the rules for rotating ETFs, how

to handle event risks, and how to overlay volatility and leverage controls. Think of this as the blueprint of the engine itself.

2.1 Identifying Macroeconomic Signals

A successful macro strategy hinges on choosing the right **signals** – quantitative indicators that provide insight into the economy's direction. Here we select a set of key macro indicators and describe how the strategy will use them:

- **Business Cycle Indicators:** The state of the business cycle (early expansion, mid-cycle, late-cycle, or recession) heavily influences which investments perform well. For example, during an early-cycle recovery, economically sensitive assets like stocks tend to rally strongly, whereas in a recession, defensive assets like Treasurys or gold shine [fidelity.com](https://www.fidelity.com). We will use indicators such as GDP growth rate, industrial production, and unemployment trends to infer the phase of the cycle. Additionally, composite indexes like the Conference Board Leading Economic Index or the PMI (Purchasing Managers' Index) can provide early warnings of cycle turns.
- **Inflation and Interest Rates:** Inflation levels (e.g. CPI from BLS) and central bank interest rate policy are crucial. Rising inflation and tightening interest rates might prompt the strategy to shift out of long-duration bonds (since their prices fall when rates rise) and into inflation-hedged assets or

commodities. Conversely, if inflation is below target and central banks are cutting rates, the strategy might increase equity and bond exposure due to a stimulative environment. Yield curve shape (spread between long-term and short-term Treasury yields) is another signal, an inverted yield curve (short rates above long rates) often signals recession, which could trigger a defensive repositioning.

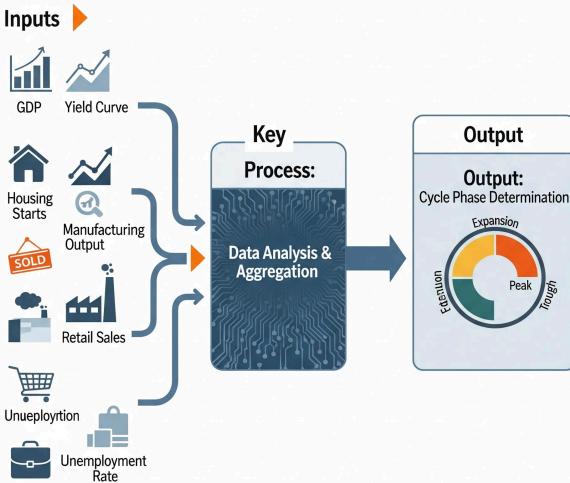
- **Credit and Financial Conditions:** We monitor credit spreads (difference in yields between corporate bonds and Treasurys) as a gauge of financial stress. Widening credit spreads can be a red flag that markets are turning risk-averse, suggesting a reduction in equity or high-yield bond ETF exposure. Other financial conditions indexes (such as the Chicago Fed National Financial Conditions Index) might be used to quantify how loose or tight overall conditions are.
- **Global Macro Inputs:** While U.S.-focused, we can't ignore global factors. Major economies' data (e.g. China's manufacturing data or European Central Bank policy) can indirectly affect U.S. markets. The strategy may include a few global signals (like USD strength, emerging market indices, etc.) if they significantly sway the ETFs in our portfolio (for instance, a strong dollar often correlates with weaker commodities and emerging markets).

Meta-Prompt Sidebar – Defining Macro Signals: Below is a pseudo-code style outline of how an LLM or algorithm would assess macro signals each cycle:

```
# Determine Macro Regime (simplified logic)
cycle_phase = None
if GDP_growth > 3% and LeadingIndicators
    rising:
        cycle_phase = "Early Expansion"
    elif GDP_growth > 0% and not all indicators
        rising:
            cycle_phase = "Mid Cycle"
    elif GDP_growth < 1% and yield_curve < 0:
        cycle_phase = "Late Cycle
        (pre-recession)"
    elif GDP_growth < 0% or recession_prob >
        50%:
        cycle_phase = "Recession"
# (In practice, use composite scores or
machine learning on many indicators)
```

This illustrates the kind of rule-based approach we'll refine in the next sections. The *cycle_phase* variable determined here will guide our sector rotation in the portfolio.

Macro Indicator Data Flow – Chapter 2.1 (Identifying Macroeconomic Signals)



*Codifies Step 1 of the strategy

Macro Indicator Data Flow Chapter 2.1 (Identifying Macroeconomic Signals) **Shows inputs → process:** Visually maps key indicators (GDP, Yield Curve, etc.) and how they flow into the single output: **Cycle Phase Determination.** This codifies Step 1 of the strategy.

2.2 Sector Rotation Across Business Cycles

Sector rotation is a classic macro strategy that we incorporate as a core of Hedge Engine. The premise is straightforward: as the economy moves through different **business cycle phases**, certain stock market sectors consistently outperform while others lag fidelity.com. By rotating into favorable sectors

(and out of unfavorable ones) at the right times, the portfolio can potentially achieve returns above the broad market.

Business Cycle Phases and Sector Performance:

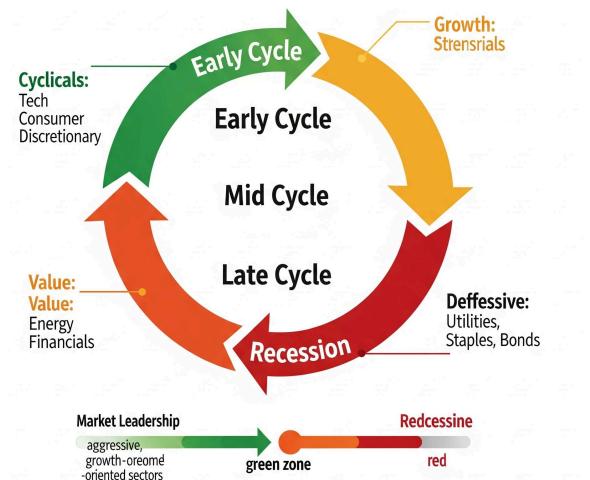
Analysts commonly break the cycle into four phases, *Early*, *Mid*, *Late*, and *Recession* each with characteristic winners and losers.

- **Early-Cycle (Recovery):** This phase follows a recession. Growth turns positive and often accelerates sharply [fidelity.com](#). Credit is easy and policy stimulus is usually in place [fidelity.com](#). **Cyclical sectors** thrive here. Historically, sectors like **Consumer Discretionary** (retail, autos), **Industrials**, and **Technology** outperform early in expansions as consumer spending picks up and businesses ramp up production. **Financials** can also do well as lending activity rises. The strategy during this phase will overweight ETFs such as technology (e.g. QQQ or XLK), industrials (XLI), consumer discretionary (XLY), and possibly small-cap stocks, which tend to rebound strongly.
- **Mid-Cycle (Expansion):** Growth is steady but not surging; the economy is typically in a more mature expansion with solid corporate profits [fidelity.com](#). Most sectors can do reasonably well here, so the focus is on **diversification**. The strategy might hold a broad market ETF (like SPY or VTI) or equal-weight sector positions. **Technology and Communication** sectors often still show strength if innovation cycles continue, but leadership can be

mixed. We might gradually trim pure cyclicals and start adding some **Quality or Defensive** tilts as the cycle ages.

- **Late-Cycle:** Signs of overheating appear, inflation may be climbing, interest rates have been rising, and growth is slowing down[fidelity.com](#). This environment favors **defensive sectors**. **Energy** and **Materials** sometimes perform decently if inflation is driven by commodity prices. **Health Care** and **Consumer Staples** (XLP) tend to outperform as consumers prioritize essentials and investors seek earnings stability. **Utilities** (XLU) can also be attractive for their stable dividends when growth is stalling. The strategy in late-cycle will underweight high-beta sectors and increase allocations to Staples, Health Care, Utilities, and possibly short-term bond ETFs to preserve capital.
- **Recession:** Economic activity contracts[fidelity.com](#) profits fall, credit is scarce[fidelity.com](#). In recessions, **preservation of capital** is key. Historically, **Treasury bonds** and **Gold** have delivered some of the best relative performance during recessions[fidelity.com](#). Equity exposure should be at a minimum, focusing only on the most defensive stocks if any (utilities, staples). The strategy might hold a Treasury bond ETF (TLT for long-term, or IEF for intermediate), increase cash or cash-like holdings, and perhaps gold or gold-miner ETFs (GLD or GDX) as a hedge. Once signs emerge that the recession is bottoming (e.g. central bank easing

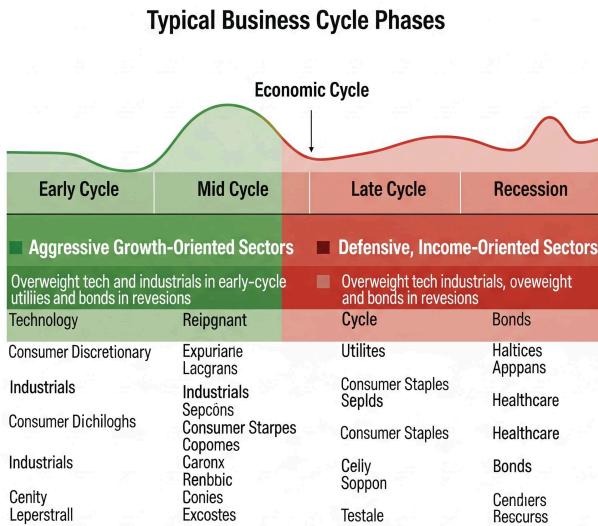
aggressively, or a fiscal stimulus), the strategy can start rotating back to early-cycle posture to capture the recovery.



An investor following a sector rotation strategy will adjust ETF allocations accordingly: e.g., overweight industrials in early-cycle, and overweight utilities and bonds in recessions.

 Fidelity.com

Typical business cycle phases and sector leadership. In early expansions, cyclicals like tech and consumer discretionary (green) tend to lead. In recessions, defensive areas (red) like utilities and staples, along with bonds, outperform.



fidelity.com

Figure: The chart above illustrates a stylized business cycle and how sector performance often rotates. As the economy moves from **Early** to **Mid** to **Late** cycle and then into **Recession**, the market leadership transitions from aggressive growth-oriented sectors (green zone) to defensive, income-oriented sectors (red zone). An investor following a sector rotation strategy will adjust ETF allocations accordingly: e.g., **overweight tech and industrials in early-cycle**, and **overweight utilities and bonds in recessions**fidelity.com.

Implementation: The Hedge Engine will use a **Sector ETF Portfolio** representing major sectors (perhaps the 11 GICS sectors via ETFs like XLK, XLF, XLY, XLP, etc., or broad

proxies for groups of sectors). Based on our determined cycle_phase (from section 2.1's signals), we'll programmatically tilt this sector portfolio. For example:

- If `cycle_phase == "Early Expansion"`: allocate higher weights to cyclical sector ETFs (Tech, Industrials, Consumer Discretionary, Financials) and lower weights to defensives (Utilities, Staples).
- If `cycle_phase == "Late Cycle"`: do the opposite – increase defensives (Staples, Health Care, Utilities) and reduce cyclicals. Possibly add inflation hedges like commodity ETFs.
- If `cycle_phase == "Recession"`: minimize equity exposure overall; hold mostly bonds (e.g. a Treasury ETF) and perhaps a small allocation to defensive equities if any.

This dynamic allocation aims to “time” the sector performance differences that come with the economic cycle. It’s important to note that **timing the cycle is challenging** – signals can be early or late. Therefore, Hedge Engine will incorporate some **smoothing** (not flipping all at once, but gradually rotating) and **diversification** to avoid all-or-nothing bets. For instance, even if our indicators strongly signal early-cycle, we might still keep some allocation to defensives in case of error, and vice versa.

Meta-Prompt Sidebar – Pseudo-Code for Sector Rotation:

```
# Pseudocode: Determine sector weights based
on cycle phase
if cycle_phase == "Early Expansion":
    weights = {"Tech":20%,
    "Industrials":15%, "ConsDiscr":15%,
    "Financials":15%,
        "Materials":10%, "Energy":5%,
    "HealthCare":5%, "ConsStaples":5%,
        "Utilities":2%,
    "RealEstate":3%, "CommServices":5%}
elif cycle_phase == "Mid Cycle":
    # Balanced allocation (approximate
market weights or equal weights)
    weights = {"Tech":15%,
    "Industrials":10%, "ConsDiscr":10%,
    "Financials":10%,
        "Materials":7%, "Energy":7%,
    "HealthCare":10%, "ConsStaples":10%,
        "Utilities":5%,
    "RealEstate":8%, "CommServices":8%}
elif cycle_phase == "Late Cycle":
    weights = {"Tech":10%, "Industrials":5%,
    "ConsDiscr":5%, "Financials":5%,
        "Materials":5%, "Energy":5%,
    "HealthCare":15%, "ConsStaples":15%,
```

```
        "Utilities":15%,  
    "RealEstate":5%, "CommServices":5%,  
    "Bonds/Cash":15%}  
elif cycle_phase == "Recession":  
    weights = {"TreasuryBondETF": 50%,  
    "GoldETF": 10%,  
        "Utilities":10%,  
    "ConsStaples":10%, "Cash":20%}  
# Note: Actual strategy might use fewer  
sectors or group some together; this is  
illustrative.
```

These weightings are examples to show the relative tilts. In practice, one would fine-tune them and possibly use optimization or historical data to decide exact weights. But the LLM (or your own implementation) can use a structure like the above to adjust the portfolio as the cycle phase changes.

2.3 Event-Driven Triggers and Probabilistic Timing

Macroeconomic indicators evolve gradually, but **market-moving events** can occur suddenly. Hedge Engine incorporates an event-driven component to handle such situations. There are two categories of events we consider:

Scheduled Events (Known Timing): These include central bank meetings (e.g., Federal Reserve's FOMC

announcements), economic data releases (monthly jobs report, inflation CPI release, GDP quarterly reports), earnings seasons, and elections or referendums. We know *when* these will happen, allowing the strategy to prepare. For each such event, we assess potential outcomes and assign probabilities. For example, before an FOMC meeting, markets might price an 80% chance of a 0.25% rate hike and 20% chance of no change. If our strategy or view differs (say we think a hike is certain), we might position accordingly – perhaps short a bond ETF or reduce equity exposure (since rate hikes can pressure stocks).

Crucially, we don't gamble outright on a single outcome; we use **probabilistic timing**. This means scaling our exposure based on the probability and impact of an event. If an event has a high probability of a certain outcome that would hurt our portfolio, we take partial protective action in proportion to that probability. If uncertainty is very high, we might temporarily reduce risky exposures across the board heading into the event. This approach helps avoid extreme positioning on low-confidence events, reducing the risk of being badly wrong.

Unscheduled or Surprise Events: These include things like geopolitical conflicts, natural disasters, sudden regulatory changes, pandemics (as 2020 taught us), or financial crises that emerge with little warning. By nature, these can't be timed, but the strategy can define triggers that attempt to detect and react to shocks:

- Use market-based sensors: e.g., if the VIX (volatility index) spikes above a threshold outside of normal

range, that signals a surge in fear and possibly an unfolding event. The strategy might then quickly de-risk (shift more into bonds or cash) even without knowing the event details.

- Stop-loss or drawdown triggers: If the portfolio starts losing beyond a set percentage in a short time (say >5% in a week outside normal volatility), it could trigger an automatic reduction in risk positions. This is a *circuit-breaker* approach to contain damage from unforeseen events.

Examples of Event Triggers and Actions:

- **Federal Reserve Rate Decision:** If the Fed is meeting and a rate hike is widely expected, the impact on markets could be negative for bonds (yields up) and potentially for stocks. Hedge Engine might *pre-emptively trim* bond ETF exposure or add an inverse bond ETF (like TBT) to hedge, and maybe lower equity exposure slightly. Once the event passes, the strategy re-evaluates – sometimes markets “sell the rumor, buy the fact,” meaning if the outcome was expected, markets might actually rebound after the announcement. The strategy’s rules could thus include reverting some of the hedge post-event if no negative surprise occurred. Conversely, if an unexpected outcome occurs (say no hike when one was expected, or a larger hike than expected), the strategy reacts at the next opportunity: e.g., if a surprise dovish move (no

hike), quickly increase equity/credit exposure to catch the upside.

- **Economic Data Release (e.g. Jobs Report or CPI):** Suppose monthly U.S. Nonfarm Payrolls (jobs) report is due. The consensus forecast is +200k jobs. We might set triggers like: if the actual figure is *much lower* than expected (signaling economic weakness), then within the day, shift allocation from cyclical equities into defensive ones or bonds, anticipating a market reaction to potential slowdown. If the figure is *much higher* (economy too hot), perhaps trim some bond exposure (since yields may rise) and consider that the Fed might hike rates faster (so maybe lighten rate-sensitive sector ETFs like real estate). These moves can be somewhat automated by defining threshold surprises (e.g., deviation from forecast) that cause trades.
- **Earnings Season and Corporate Events:** Although this is more micro than macro, broad-based earnings trends (e.g., if in a given quarter, a majority of companies are surprising positively vs. estimates) can shape the macro outlook. The strategy might include a check during quarterly earnings season: if earnings are robust and guidance strong across sectors, it reinforces an optimistic macro stance (maybe stay invested in equities despite other worries). If earnings are crumbling, it may strengthen a bearish macro signal.

- **Geopolitical Crisis:** If news breaks of a geopolitical conflict, the immediate response rule might be: increase gold and Treasury bond exposure (traditional safe havens) and reduce equities, especially those with high international exposure or those directly impacted (e.g., an oil supply shock event would cause a spike in oil – so energy ETF might actually go *up*, while broad market falls). These decisions can be hard-coded to an extent (e.g., “if conflict in Middle East, then Energy ETF upweight, Airlines ETF downweight, etc.”) but generally, an LLM or human oversight is needed to qualitatively assess the situation. The strategy may simply have a contingency: “In case of major crisis headlines, temporarily move X% to cash and safe assets, and await clarity.”

Probabilistic Position Sizing: A key principle is that the strategy’s reaction size is proportional to both the **confidence** in the event outcome and the **estimated impact** of that outcome. For instance, if there’s a 60% chance of an event that could cause a 10% market drop, one might position for roughly a 6% downside ($0.6 * 10\%$). In practice, this could mean hedging 60% of the portfolio’s equity exposure. If the probability rises (new info comes in), the hedge can be increased. This dynamic sizing prevents the strategy from being all-in on low-probability bets, but still provides protection or opportunistic positioning where risk/reward is favorable.

Meta-Prompt Sidebar – Event Trigger Pseudocode:

```

# Pseudocode for event handling (simplified
example for a Fed meeting)
if days_to_next_FOMC <= 5:
    prob_hike =
market_implied_probability("rate_hike") #  

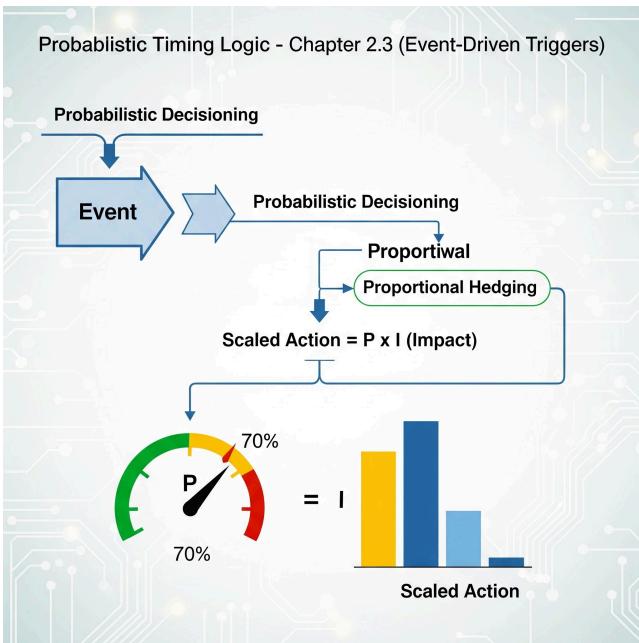
e.g. from Fed futures
    if prob_hike > 0.7:
        # Likely hike: hedge interest-rate
sensitive assets
            decrease_exposure("TreasuryBondETF",
amount=prob_hike * base_reduction)
            decrease_exposure("REITs",
amount=prob_hike * base_reduction) #  

rate-sensitive sector
    elif prob_hike < 0.3:
        # Likely no hike or cut: position
for rally in bonds/stocks
            increase_exposure("TreasuryBondETF",
amount=(1-prob_hike) * base_increase)
            increase_exposure("TechSectorETF",
amount=(1-prob_hike) * base_increase)
# After event, adjust based on outcome:
if event_passed("FOMC"):
    outcome = get_FOMC_outcome() # e.g.
"hike" or "no_change"
    if outcome == "hike" and prob_hike <
0.5:

```

```
    # surprise hawkish, markets likely
    drop -> go defensive
        shift_portfolio("defensive")
    elif outcome == "no_change" and
    prob_hike > 0.5:
        # surprise dovish -> increase risk
        exposure
            shift_portfolio("risk_on")
```

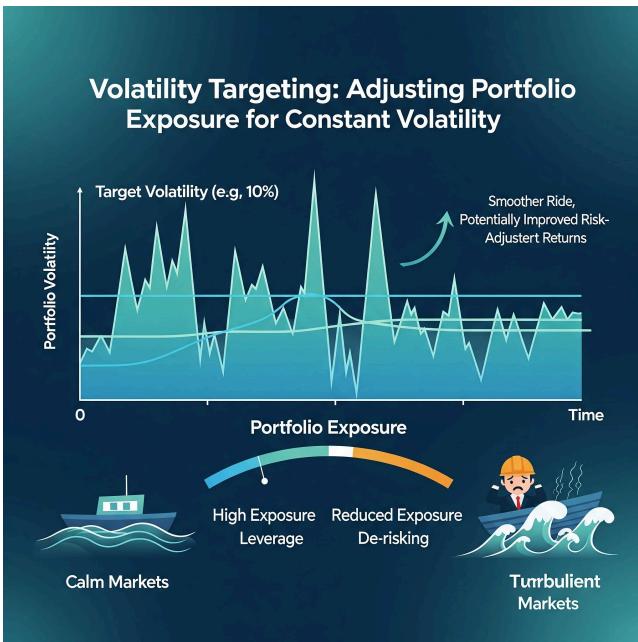
This pseudo-code outlines a framework: pre-event hedging based on probabilities, followed by post-event adjustment based on surprise vs expectation. An LLM using the eBook as context could populate such a template with real data and take action accordingly.



Probabilistic Timing Logic Chapter 2.3 (Event-Driven Triggers) **Shows probabilistic decisioning:** Clarifies the concept of proportional hedging by illustrating how a probability P and an impact I result in a scaled action ($\$P \backslash \text{times } I \$$). This is crucial for the Event Triggers.

2.4 Volatility Targeting and Synthetic Leverage

A distinguishing feature of the Hedge Engine strategy is its use of **volatility targeting** and **synthetic leverage** to manage risk and enhance returns. These are techniques often used by institutional funds:



How we implement it: The strategy will regularly calculate the portfolio's **realized volatility** (say using a 20-day or 1-month rolling window of returns). If the realized vol is below the desired target, we will *increase* exposure (either by allocating more to high-volatility assets or by using leverage) until expected volatility is back near target. If realized vol is above target (market swinging more than we want), we *decrease* exposure (shift some assets to cash or safer assets) to bring volatility down. This can be done by scaling all positions up or down multiplicatively.

Consider an example: The target volatility is 10% annual (~approx 0.63% daily). If the portfolio's recent volatility has been only 5% annual (very quiet market), the strategy might

double its exposure (i.e., use 2:1 leverage) to reach the 10% risk target. Conversely, if volatility jumps to 20% annual, the strategy might halve its exposure to protect against large swings. This dynamic scaling often results in **buying low volatility and selling high volatility** – which tends to reduce drawdowns. Indeed, research has shown that volatility-managed portfolios can sidestep some of the worst drawdowns of unmanaged portfolios by cutting risk in stress periods researchaffiliates.com.

Synthetic Leverage and Overlays: When we say synthetic leverage, we refer to using instruments like **leveraged ETFs, futures, or options** to increase exposure without having to put in dollar-for-dollar capital. For example, instead of buying \$100 of an index, one could buy \$50 of a 2x leveraged ETF of that index to get roughly the same exposure. The Hedge Engine might use such tools *sparingly* and systematically. One use case: if our signals are strongly bullish and volatility is low, the strategy could employ a **leverage overlay** to boost returns, for instance, using a 1.5× exposure to equities temporarily. Conversely, in certain hedging situations, an *inverse* or leveraged inverse ETF can quickly reduce net exposure (e.g., buying a -2x S&P500 ETF as a hedge during a crash). We call it “synthetic” because it doesn’t require borrowing cash on margin in the traditional sense; the leverage is embedded in the instruments.

It’s critical to highlight the **risks of leverage**: leveraged ETFs and derivatives can amplify losses and often have quirks (leveraged ETFs typically aim for daily percentage goals and can drift from their target over longer periods due

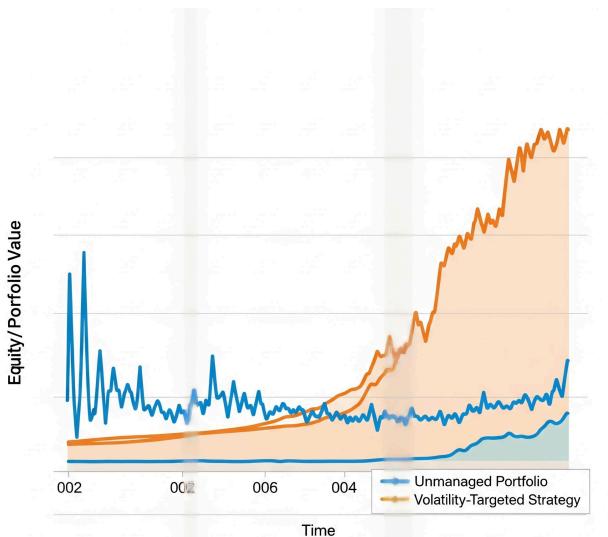
to compounding effects). These products are **considered risky and often suitable only for short holding periods**, as noted in brokerage risk disclosures [tiaa.org](#). For example, a 2x ETF held over a volatile month may not exactly double the index's monthly return, it could underperform due to volatility decay. The strategy must account for this, potentially by resetting leveraged positions frequently or only using them short-term around specific opportunities.

Volatility Overlay: Another related concept is a **volatility trading overlay** using volatility instruments (like VIX futures or options) either to hedge or enhance the portfolio. For instance, if volatility spikes (VIX goes very high, indicating fear), the strategy might take a small long position in VIX-related ETF (e.g., VXX) as a hedge against further turmoil or to profit from mean reversion (as volatility usually falls from extremes). Conversely, during very calm periods, a *cautious* short volatility position (via inverse VIX ETF or selling options) might generate extra income, but this is an advanced move and can be dangerous (short volatility strategies can blow up in a sudden crisis). We will keep any vol overlay component limited and always tied to the volatility targeting framework (i.e., it should not increase overall portfolio risk beyond the target).

Let's illustrate the effect of volatility targeting with a hypothetical comparison. Consider two portfolios: one static (unmanaged) and one volatility-targeted. Suppose both start with \$100 and invest in the same asset, but the vol-targeted one scales exposure inversely with volatility.

Above: Simulated equity curves for an unmanaged portfolio (blue) versus a volatility-targeted strategy (orange) over a sample period. The volatility-targeted strategy adjusts exposure to keep risk steady, resulting in a smoother equity curve with fewer steep drops. In calm periods, the orange line rises faster (because the strategy took on more exposure), while in volatile periods, it flattens (having cut exposure), thus avoiding some losses. Over time, the volatility-targeted approach can achieve similar or slightly higher returns with reduced drawdowns.

The simulated equity curves clearly demonstrate the benefit of the volatility-targeted strategy (orange) compared to an unmanaged portfolio (blue).



Above: Drawdown (%) comparison for the same two strategies. The unmanaged portfolio (blue) experiences deeper drawdowns (worst declines reaching around -30% in this example), whereas the volatility-targeted portfolio (orange) limits drawdowns to smaller magnitudes. By scaling down risk in turbulent times, the vol-targeted strategy avoids the full brunt of market crashes. This aligns with research findings that volatility targeting can effectively manage drawdowns and volatility more effectively than an unmanaged approach researchaffiliates.com. However, note that if the market experiences prolonged low-volatility rallies, the vol-target strategy might underperform slightly (since it would be leveraged up and could suffer if a sudden spike catches it). Thus, risk controls on leverage are still essential.

Practical notes on implementation: We will set a **volatility measurement window** (e.g., 20 trading days) and a target vol (e.g., 10% annual). Each day or week, the strategy computes recent vol. It then computes a scaling factor = (Target Vol / Recent Vol). For example, if recent vol = 5%, Target = 10%, scale = $10/5 = 2\times$, meaning double exposure. We also impose limits (perhaps never go above $2\times$ leverage or below, say, $0.3\times$ exposure for extremely volatile times, to ensure we always keep some skin in the game but never over-lever too far). This scaling factor is applied to all risk assets in the portfolio proportionally. If using leveraged ETFs, we might explicitly shift holdings to those rather than borrowing on margin. For instance, if we

want 1.5× exposure to stocks, we might mix in some 2x S&P ETF (like SSO) into the equity allocation.

From a meta-prompt perspective, the LLM reading this can extract a clear algorithm: measure volatility → adjust exposure. We can codify it for clarity:

Meta-Prompt Sidebar – Volatility Targeting Algorithm:

```
# Volatility Targeting Adjustment (run
periodically, e.g., weekly)
target_vol = 0.10 # 10% annual target
recent_vol =
calculate_annualized_vol(portfolio_returns_1
ast_20d)
scale_factor = target_vol / recent_vol

# Limit the scale factor to avoid extreme
leverage:
scale_factor = min(scale_factor, 2.0)
scale_factor = max(scale_factor, 0.5)

apply_scale_to_portfolio(scale_factor)
# (If scale_factor > 1, and margin is
available, use leveraged ETFs or futures to
increase exposure.
# If scale_factor < 1, move a portion of
portfolio to cash or bonds to reduce
exposure.)
```

This pseudo-code encapsulates the vol targeting rule. The LLM can use such logic to maintain the desired risk profile. All of this, of course, operates within the broader strategy – e.g., it will scale the *sector-rotated* portfolio we determined earlier, rather than altering the composition.

2.5 Synthesis of Strategy Rules

To wrap up the design: the Hedge Engine strategy brings together **macro signals**, **sector rotation**, **event triggers**, and **vol targeting** into one coherent system. In practice, each trading period (say each week or month), the strategy would:

1. **Update Macroeconomic Data:** Pull the latest values for key indicators (GDP nowcasts, CPI, unemployment, PMI, yield curves, etc.) from reliable sources like FRED or BLS. Update any probability models (e.g., recession probability or Fed hike probability).
2. **Determine Regime & Allocation:** Based on those indicators, decide the current macro regime (expansion, recession, etc.) and set the target sector/allocation weights accordingly (per section 2.2 rules).
3. **Incorporate Event Outlook:** Check the calendar for upcoming events (next 1-2 weeks). Adjust the allocation or hedges for any high-impact events

using the probabilistic rules (from 2.3). For example, if a major event is imminent, perhaps hold a bit more cash or a put option as hedge.

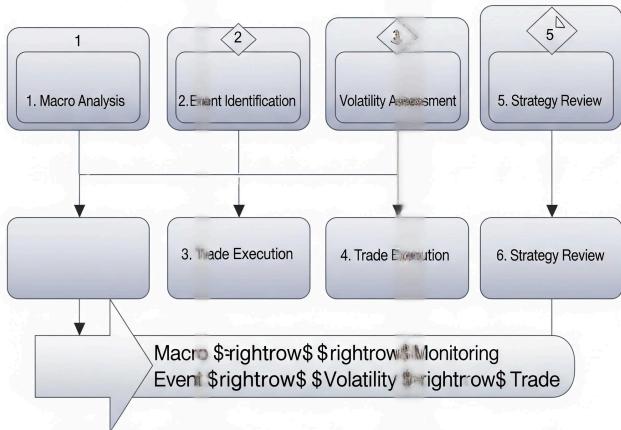
4. **Apply Volatility Scaling:** Compute the current portfolio volatility and apply the scaling factor (from 2.4). This changes the overall exposure level (and may involve switching some portion to leveraged instruments if scaling >1 , or to cash if scaling <1).
5. **Trade Execution:** Implement the changes through trades in the chosen ETFs. This could mean rotating weights (buy some, sell others) and setting up hedges. The strategy likely trades at a set frequency (e.g., monthly rotation with allowances for intra-month adjustments around big events).
6. **Monitoring & Stops:** Continuously monitor portfolio performance and risk. If intra-month something breaches a risk limit (like portfolio down X% from peak, or a volatility jump), the strategy might override normal schedule and rebalance immediately to cut risk (like an emergency brake).

Each of these steps can be automated or followed by an AI agent if it's provided with the right data and this rulebook (the meta-prompt). Next, we will discuss how to evaluate and backtest this strategy to ensure it has merit, and then how to implement it in real life, including data sources and legal considerations.

Six-Step Strategy Execution Flowchart –

Chapter 2.5 (Synthesis of Strategy Rules)

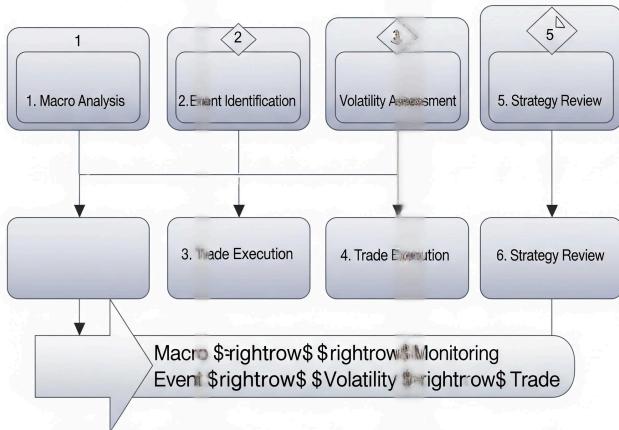
Shows the full algorithm: Provides a high-level map of the Hedge Engine. It explicitly shows the required sequence: Macro → Event → Volatility → rightrows as the definitive LLM instruction flow



Six-Step Strategy Execution Flowchart –

Chapter 2.5 (Synthesis of Strategy Rules)

Shows the full algorithm: Provides a high-level map of the Hedge Engine. It explicitly shows the required sequence: Macro → Event → Volatility → rightrows as the definitive LLM instruction flow.



3. Backtesting and Performance Evaluation

Designing a strategy is only half the battle – we need to **backtest** it on historical data to gauge how it would have performed, and to refine the rules. Backtesting a macro ETF strategy comes with challenges, but also huge benefits: it can reveal how the strategy might handle past recessions, bull markets, and shocks, and whether its risk management holds up.

3.1 Backtest Framework:

To backtest Hedge Engine, you'd follow these steps:

- **Data Collection:** Gather historical time series for all inputs:
 - Economic indicators (e.g., monthly GDP growth estimates, inflation rates, etc. – FRED[reddit.com](#) is a great source for many U.S. macro series, and BLS[reddit.com](#) provides core data like CPI and unemployment).
 - ETF price histories for all the ETFs in the strategy (sector ETFs, broad market ETFs, bond ETFs, etc.). Sources like Yahoo Finance or investing APIs can provide daily price data. Ensure to include dividends for total return if significant.
 - Event dates and outcomes: create a calendar of past Fed meetings, major data releases, etc., along with what happened (for conditional event logic testing).
- **Simulation Engine:** Step through time (likely monthly or weekly increments). At each step, feed the historical indicator values into the strategy logic exactly as if we were at that point in time.
Important: Use only data that would have been known at that time (to avoid look-ahead bias). For example, if using GDP, note that GDP is reported with a lag; a sophisticated backtest might use the data as of the release dates. (Advanced: use “vintage” datasets or ALFRED for FRED to get data

as reported at the time, since macro data often gets revised [reddit.com](#). However, for a rough test, using final data may be a start, but be cautious.)

- **Apply Trades:** Simulate the rebalancing of the portfolio according to strategy rules: rotating sector weights, scaling exposure, etc. Account for transaction costs (ETFs have small spreads, but if rotating monthly it's minimal) and any leverage financing costs if applicable.
- **Record Performance:** Track the portfolio value over time. Also track key metrics: returns, volatility, max drawdown, Sharpe ratio, Calmar ratio (return/drawdown), etc. Compare to benchmarks: e.g., how did Hedge Engine do vs a simple 60/40 stock-bond portfolio, or vs S&P 500, or vs a global macro hedge fund index if available.

3.2 Interpreting Results:

You would examine whether the strategy indeed provided smoother performance, reduced drawdowns, or enhanced returns:

- Did the **sector rotation** add value? For instance, during known recessions (2001, 2008, 2020), did the strategy successfully tilt defensively and avoid big losses relative to the market?
- Did **volatility targeting** reduce the drawdown in 2008 or March 2020? Often a vol-target strategy

would have de-levered before or as volatility spiked, thereby experiencing less of the crash.

- How did the strategy perform in long bull markets (e.g., 2010s)? Vol targeting might cause underperformance in steadily rising low-vol markets (because it would leverage up and then possibly get caught by sudden corrections, or simply because being unlevered buy-and-hold might win if there were no major vol spikes to dodge). Check if the strategy still kept up, or if not, whether risk-adjusted returns were better.
- **Drawdowns & Recovery:** Look at the worst peak-to-trough declines (drawdowns). The chart above illustrated how vol targeting can mitigate drawdowns. Verify this in the backtest: e.g., if S&P 500 had -50% in 2008, maybe Hedge Engine was down much less thanks to rotation into bonds and cutting exposure. Also see how quickly it recovers from drawdowns – a good strategy should recover faster by getting back into risk at the right times.
- **Performance windows:** A concept in macro evaluation is to check performance in different environments or “windows” – e.g., performance in recession vs expansion periods, performance in high-volatility vs low-volatility regimes, performance specifically around events (did it handle the 2016 Brexit event or the 2020 COVID crash effectively?). This can be visualized by plotting returns conditioned on some factor or simply by annotating

the equity curve with regime labels.

Additionally, **visualizations are invaluable** in explaining and validating the strategy:

- **Equity Curve with Benchmarks:** Plot the cumulative return of Hedge Engine vs S&P 500 and maybe vs a 60/40 portfolio. This shows at a glance if the strategy is adding value. Ideally, the Hedge Engine line would be smoother and higher at the end than the benchmarks (though even if final value is similar, a smoother ride is a win for many investors).
- **Drawdown Chart:** We included a drawdown comparison figure (above) – that kind of chart explicitly shows every drawdown over time. Perhaps Hedge Engine's worst drawdown was, say, -15% while S&P's was -50%. Seeing the magnitude and duration of drawdowns helps investors gauge risk (e.g., "if I invest in this, the worst I might expect is a 15% loss at some point").
- **Rolling Risk/Return Metrics:** One can plot rolling 3-year Sharpe ratios or volatility to show consistency. A strategy that is truly stable will show relatively stable volatility over time (since that's a goal here).
- **Attribution Charts:** For deeper analysis, break down where returns are coming from – e.g., how

much of the outperformance (if any) came from sector selection vs. from volatility timing vs. from event trades. This might involve comparing to variants of the strategy (like a version without vol targeting, or without event triggers) to see each component's value-add.

Backtest Pitfalls: It's important to not overfit the strategy to historical data. Because we are using a variety of well-known principles (cycle rotation, etc.), we're less likely to over-optimize specific parameters. Still, one should test variations (sensitivity analysis). For example, if you changed the threshold for an event trigger or the target volatility level, does the performance dramatically change? If a strategy only works with very precise parameter values, it might be curve-fit and not robust. We prefer robust, broadly effective settings.

In evaluating backtest results, keep a bit of skepticism:

- **Look for Regime Robustness:** Does the strategy hold up in different decades? Macro relationships can change (e.g., in some cycles tech might not boom as expected, or in some recessions policy responses can make stocks rally earlier).
- **Include Realistic Frictions:** If the strategy trades very often around events, consider trading costs and potential slippage. Our design is moderate frequency (monthly with occasional tactical moves),

so it should be fine for ETFs.

- **Beware of “lookahead” biases:** Ensure your backtest truly simulated information flow correctly. For instance, if using an indicator that is revised later, the backtest should use initial releases, or else it may paint an unrealistically rosy picture (using data that wasn't actually known at the time)[reddit.com](https://www.reddit.com).

A well-executed backtest, coupled with **out-of-sample tests** (like testing on a period not used in strategy development, or on international data for example) can give confidence that Hedge Engine's methodology is sound. Assuming the results are favorable – showing improved risk-adjusted returns, smaller drawdowns, etc. – we can proceed to consider implementation details. Before real-world rollout though, we must address another critical aspect: **LLM integration** and how to format our strategy as something an AI can directly use.

4. LLM Integration: Crafting the Strategy as a Meta-Prompt

One innovative aspect of this eBook is that it doubles as a “program” for an AI. We want an LLM to effectively ingest this strategy description and then be able to *act as a hedge fund manager*, following the rules to make decisions. To achieve that, we pay special attention to prompt engineering in the way the content is written.

4.1 Clear Structure and Tags: We have structured each section with clear headings and even used pseudo-code blocks. An LLM can use the headings to find relevant parts (e.g., if it needs to recall how to do volatility targeting, it can search within the text for “Volatility Targeting Algorithm” or find Section 2.4). By delineating the logic in step-by-step lists and code-like format, we reduce ambiguity. For example, the pseudo-code in earlier sections can be directly converted to actual code with minor effort. If the LLM is asked to “generate the current portfolio allocation given these macro inputs,” it can follow the outlined steps: identify cycle phase, set weights, check events, apply vol scale.

4.2 Meta-Prompt Sidebars: Throughout the chapters, we included **Meta-Prompt Sidebar** content. These sidebars are essentially instructions or code fragments intended for the LLM’s benefit. They distill the preceding explanatory text into actionable logic. In a final AI prompt, one might instruct the LLM to **extract all pseudo-code blocks** from the Hedge Engine document and execute them in order, using current data. We’ve made those blocks self-contained enough that this is feasible. For instance, the sidebar in 2.2 gave a mapping from cycle phase to sector weights – an LLM could store that as a dictionary in memory. The event trigger pseudo-code was more conditional, but again, structured clearly with **if...elif** logic that the LLM can interpret.

A meta-prompt approach might involve first prompting the LLM with this entire eBook as context, then giving it a task like: *“Using the methodology described in Hedge Engine,*

analyze the following current data and output the recommended portfolio trades." Because the content is presented systematically, the LLM doesn't have to infer rules from narrative – we have spelled them out. This dramatically lowers the chance of misinterpretation or "hallucination" of strategies not intended. Essentially, we've given the LLM an internal knowledge base in the form of this well-structured guide.

4.3 Prompt Blueprint Example: Below is an example of how one could construct a prompt for the LLM after feeding it this guide. Note that this is more for the human implementer to see how the pieces come together:

[System message to LLM:]

You are an expert macro portfolio AI following the Hedge Engine strategy described in the provided document. Your role is to read real-time macroeconomic and market data, then apply the strategy rules to manage a portfolio of ETFs.

[User message:]

Here is the current market data (as of YYYY-MM-DD):

- GDP Nowcast growth: 1.2% (slowing)
- CPI year-over-year: 4.1% (above target)
- Yield curve (10y-2y spread): -0.25% (inverted)
- Unemployment rate: 4.8% (rising)

- Fed Funds Rate: 5.00%, next FOMC in 10 days, probability of +25bp hike = 60%
- VIX: 18
- Portfolio current allocation: {Tech 15%, Industrials 5%, ... etc.}

Using the Hedge Engine methodology, determine:

1. The current business cycle phase and any allocation shifts (sector rotation) needed.
2. Any event-driven adjustments given the upcoming FOMC.
3. The volatility scaling factor and whether to lever or de-lever the portfolio.
4. The final portfolio allocation and list of trades to get there.

The LLM, having the context of this eBook, would internally follow something like:

- Find where cycle phase determination is described (Section 2.1/2.2) and apply it: likely we're late-cycle or approaching recession given inverted yield curve and rising unemployment. So it decides on a more defensive allocation (maybe follows that pseudo-code mapping weights to defensives).

- See that an FOMC is coming with 60% hike probability – from Section 2.3 logic, it might modestly hedge rate-sensitive assets (maybe reduce bonds a bit).
- Check volatility: VIX 18 relative to our target vol; 18 isn't extremely high, but maybe a bit above long-term average. It calculates realized vol (not given in prompt, it may assume something). Possibly slight de-risk if vol above target.
- Then output an answer with the reasoning and trades (like: “Identified cycle phase as Late Cycle -> shifting 5% from Tech to Staples... With Fed hike risk at 60%, reducing bond ETF allocation by X... Current volatility near target, scaling factor ~1, so no leverage change... Final allocation: etc.”).

4.4 Format for LLM Parsing: In the eBook text, we have been mindful to use consistent terminology and units (e.g., always specifying percentages, labeling phases exactly). We avoid ambiguous language. For instance, instead of saying “increase a bit”, we gave explicit examples like “increase exposure by 10%” or “scale factor 1.5x”. This helps an AI quantitatively interpret the guidance. Additionally, our use of markdown lists and sub-lists creates a hierarchy that an AI can traverse. Headings (##, ###) are clear boundaries of topics.

One could also instruct an LLM to only pay attention to pseudo-code blocks if needed (filtering out prose). Because

we've ensured that all crucial decision logic is mirrored in either a list or code block, doing so would not lose the strategy essence. This is by design – **the strategy is described twice: once in prose (for understanding) and once in structured form (for execution)**. This redundancy is intentional to serve the dual audience (human and AI).

4.5 Continual Learning and Adaptation: An LLM could also be tasked with reading new research or data and updating the strategy. For instance, if some new evidence suggested a different indicator is useful, one could append a note to this document or provide an updated prompt. The structured nature of this guide means new rules can be slotted in (perhaps as an added pseudo-code rule or an update to weights). For example, if we wanted to incorporate a **machine learning prediction** for recession probability, we could add a sidebar with that model's usage, and the LLM would incorporate it in its decision flow.

In summary, by writing this eBook in a **pseudo-code augmented, step-by-step format**, we've essentially created a **blueprint prompt** that an LLM can transform into an algorithm. Next, we will address the practical matters of executing this strategy in the real world – from data sourcing to legal considerations and a suggested rollout plan.

5. Risk Management and Disclosure

No investment strategy is complete without a thorough discussion of risks and proper disclosures, especially for one employing leverage, derivatives, and algorithmic

decision-making. This section serves as both guidance for managing risks within the strategy and as a template for the kind of **disclosures** that should accompany a professional strategy document.

5.1 Key Risk Factors:

- **Market Risk:** All ETFs and investments involved are subject to general market risk, the possibility of losing value due to market fluctuations. The strategy can (and will at times) suffer losses. Even with hedging and rotation, losses can occur if multiple positions move adversely or if macro signals prove wrong. **There is no guarantee of profit; one can lose principal.** For example, an unforeseen event could cause equities, bonds, and even gold to all drop together in the short term, defeating the strategy's intent to hedge via decorrelation.
- **Model/Strategy Risk:** Hedge Engine's decisions rely on models and historical relationships (business cycle patterns, etc.). These relationships may break down. For instance, if a future recession doesn't see defensives outperform (perhaps due to unique circumstances), the strategy could be caught wrong-footed. The use of an LLM to execute rules also introduces risk, the AI might misinterpret a rule or there could be an error in data input. *Any algorithmic strategy can fail if the underlying assumptions or data are flawed.*

- **Leverage Risk:** The strategy's use of leveraged ETFs or synthetic leverage magnifies gains *and losses*. A 2× leveraged position will lose about 2% if the underlying index drops 1% (minus some compounding effects). In extreme volatility, leveraged products can lose a large portion of their value rapidly gdcdyn.interactivebrokers.com. Additionally, leveraged ETFs often aim for daily targets; holding them over long periods can lead to unexpected results (they can drift from the intended multiple, especially in see-saw markets). **Investors must understand that leverage can lead to rapid, significant losses**, potentially more than the initial investment if using margin.
- **Complex Instrument Risk:** Some of the ETFs possibly used (inverse ETFs, volatility ETFs, etc.) are considered “complex products.” They may have unique risks: inverse ETFs reset daily and can behave counterintuitively if held long. Volatility ETNs/ETFs can decay in value during calm periods due to contango in futures. These instruments often come with warnings that they are for short-term tactical use only tiaa.org. Using them requires careful monitoring (which Hedge Engine’s rules do, by frequently reassessing volatility and trends).
- **Event Risk:** While we incorporate events, there is still the risk of a gap move – when markets reopen far from previous prices due to overnight news. Our stop-loss measures might not trigger in time if an event causes, say, a sudden 20% drop at the open

(e.g., a surprise announcement over a weekend). Options can be used to mitigate some gap risk (because they provide insurance), but not without cost.

- **Liquidity Risk:** Most ETFs we plan to use (major index and sector ETFs) are highly liquid. However, in a crisis, even ETFs can suffer widened bid-ask spreads or deviances from NAV. If the strategy needs to rebalance during such a time, it could face slippage. For example, bond ETFs in the March 2020 COVID crash saw liquidity issues. The event-driven nature means we might be trading when everyone else is (around events), which can be a liquidity crunch time.
- **Data and Technology Risk:** Relying on real-time data feeds (from FRED, BLS, etc.) and an LLM or code means we have operational risk. Data might be delayed or incorrect, or the API might fail at a critical moment. The AI or algorithm might crash or encounter an unforeseen situation not coded in the rules. It's crucial to have **fail-safes**: e.g., if the system is unsure, perhaps default to a safe mode (like move to a neutral 50% stock, 50% bond allocation or full cash until human oversight).
- **Regime Change Risk:** The macro regime itself could change in nature. For instance, if we enter a period of stagflation (high inflation, low growth) that doesn't fit well into the historical patterns we used (which assume either growth or recession, but not

both inflation and stagnation), the strategy might need new rules. Similarly, unprecedented monetary policy (like negative interest rates or massive QE) could make historical analogs less useful. Continuous research and adaptability are needed.

5.2 Disclosures (for a Professional Document):

A formal document for investors would include disclosures such as:

- “**Not an Offer:** This document is for informational purposes and does not constitute an offer to sell or a solicitation to buy any security or investment product.”
- “**No Personal Advice:** The strategy described is general in nature and does not take into account individual circumstances or objectives. Investors should consult with a financial advisor before implementing any strategy.”
- “**Past Performance Not Indicative:** Hypothetical backtested performance (if any is presented) is not indicative of future results. There is no assurance the strategy will achieve similar results or avoid losses.”
- “**Risk of Loss:** All investments carry risk of loss of some or all capital. This strategy involves certain additional risks, including leverage and derivatives use, which can amplify losses. Investors should only

risk capital they can afford to lose.”

- “**Complex Products:** *The use of leveraged and inverse ETFs, as well as volatility-linked instruments, introduces unique risks. Such products may not be suitable for all investors and often are designed for short-term holding periods [tiaa.org](#). One should thoroughly understand these before use. The value of leveraged ETFs can deteriorate over time in volatile markets even if the underlying index is flat.”*
- “**Algorithmic Decision-Making:** *The methodology may be executed by an algorithm or AI. While this can remove emotion and enforce discipline, it also means mechanical failures or programming errors can occur. All automated decisions should be monitored by a human for reasonableness. The strategy should not be ‘left on auto-pilot’ without oversight.”*
- “**Data Sources and Reliability:** *The strategy relies on third-party data (e.g., economic indicators from government sources). We believe these to be reliable, but we cannot guarantee their accuracy or timeliness. Data revisions are common; the strategy’s signals are only as good as the data available at the time.”*
- “**Regulatory Considerations:** *Use of leverage and frequent trading can have tax implications and may not be permitted in certain account types. Ensure compliance with all regulations (such as trading*

restrictions in retirement accounts, or any pattern day trading rules if applicable).”

- “**Forward-Looking Statements:** Any opinions or forecasts in this document are forward-looking statements. They involve risks, uncertainties, and assumptions. Actual outcomes and results could differ materially.”

Including such disclosures is standard in professional reports, especially ones outlining an investment methodology. For Hedge Engine, it's particularly important to highlight the leveraged ETF risks (e.g., how a 3× ETF can theoretically go to zero if the index falls ~33% in a day, which has happened in extreme cases).

5.3 Risk Management Practices:

Beyond disclosures, the strategy itself integrates risk management:

- We've built in **volatility control** to automatically reduce risk in turbulent times (that's a risk management device).
- We propose **position limits** (e.g., no single sector ETF > 25% of portfolio except safe assets, to ensure diversification; no use of >50% of portfolio in leveraged products, etc.).
- **Stop-loss rules:** Though not explicitly detailed earlier, one might add: if portfolio drawdown

exceeds, say, 10%, then cut all positions by X% to stop the bleed, essentially moving to protection mode.

- **Review and oversight:** At least monthly, if not more, a human (or the AI with a separate analytical routine) should review if the strategy's assumptions still hold. For example, if an indicator that usually works has started giving false signals, maybe remove or replace it. The dynamic nature of macro investing means we must adapt.

In sum, risk management isn't just a paragraph in the book it's woven through our strategy design (from diversification to volatility targeting to event hedging). Any user of this strategy or any LLM implementing it should strictly adhere to these risk protocols and fully understand the disclaimers before live trading.

6. Implementation and Rollout Plan

Designing the strategy and even testing it is still a theoretical exercise until we actually implement it with real money. In this final chapter, we outline a practical **rollout plan** for bringing Hedge Engine to life, including setting up data feeds, leveraging technology (possibly the LLM itself) for execution, and the step-by-step process from paper to practice. We'll also list some top real-time data sources to power the engine.

6.1 Phase 1: Preparation and Infrastructure

Data Sourcing: We need reliable, up-to-date feeds for:

- *Macroeconomic Data:* **FRED (Federal Reserve Economic Data)** is a prime source for U.S. economic time series, aggregating over 800k data series from 118 sources fred.stlouisfed.org. It includes GDP, CPI, employment, interest rates, etc. We can use FRED's API for programmatic access (with an API key, one can pull series like CPIAUCSL for CPI, GDPC1 for real GDP, etc.). For more granular or real-time data: the **Bureau of Labor Statistics (BLS)** provides monthly CPI and employment data on their website (and an API) reddit.com, often releasing at preset times (e.g., Jobs report first Friday of each month at 8:30am ET). The **Bureau of Economic Analysis (BEA)** provides GDP and PCE inflation data (quarterly and monthly). **Treasury** yields can be fetched via U.S. Treasury or also via FRED (they have daily yield curve series).
- *Market Data:* Price data for ETFs will likely come from financial APIs. Options include:
 - **Yahoo Finance API** (unofficial through yfinance in Python) – free daily data for most ETFs.

- **Alpha Vantage** (offers a free API key with certain limits).
 - **IEX Cloud or Tiingo** – affordable APIs for daily prices.
 - If using a trading platform like Interactive Brokers, their API could stream prices.
Ensure we have OHLCV (open-high-low-close-volume) or at least daily close and maybe real-time if doing intraday (though our strategy is not high-frequency, end-of-day or even end-of-week data might suffice).
- *Event Information:* A calendar of events can be manually maintained or pulled from sources:
 - **EconoDay or Forex Factory** provide calendars of econ releases with forecasts.
 - For Fed meetings, one can scrape the schedule from the Fed website or rely on known dates. Fed funds futures for probability can be obtained from CME's website (they publish probabilities).
 - **News feeds:** It's tricky to automate reactions to unscheduled news, but having a news alert system (RSS feeds from Bloomberg, Reuters headlines) might be helpful. An LLM connected to news could even parse

headlines for certain keywords (though caution with that to avoid false signals).

Technology Stack: Decide how to execute the strategy:

- A common approach is to use **Python** for orchestrating data and decisions, given the plethora of finance libraries. One could code the rules (some we pseudo-coded) using a library like pandas for data, and connect to a broker API for trades.
- If the LLM itself is to be used in execution (as a sort of AI portfolio manager), one way is to integrate it via an API (like OpenAI's API) and prompt it with the necessary context (the strategy and current data) periodically. However, relying on an LLM in real-time requires careful handling of its outputs (and costs, latency). It might be more straightforward to codify the strategy in a deterministic way for execution and perhaps use the LLM in a research or monitoring capacity (e.g., ask it to interpret unusual situations or summarize news impacts which quantitative code struggles with).
- **Backtesting platform:** If not done already, use something like **QuantConnect**, **Backtrader**, **Zipline**, or **Quantopian (was)** to backtest. QuantConnect, for example, has data and allows coding algorithms (and they include FRED data integrationquantconnect.com). This could accelerate

development.

- **Storage:** maintain a database or at least a CSV log of macro data and signals as they come in, for transparency and analysis.

Initial Parameter Finalization: Before going live, finalize key parameters:

- Target volatility level (maybe 10% as discussed, but ensure that's acceptable to the investor).
- Rebalancing frequency (maybe monthly base with weekly checks).
- The universe of ETFs to be used (list them clearly, e.g., SPY for broad equity, sector SPDRs for sectors, TLT for bonds, GLD for gold, SHY or BIL for cash proxy, etc., including any leveraged ones like SSO or others – ensure they're available and liquid).
- Position limits and max leverage boundaries.

6.2 Phase 2: Paper Trading and Calibration

Never jump in with full capital. The next step is **paper trading** (or using a small amount of capital as a pilot):

- Run the strategy live in simulation for a few months. Many brokerages offer paper trading accounts

(Interactive Brokers, Thinkorswim, etc.) where you can test your automated strategy as if real. Or simply simulate with daily data updates and record trades you *would* do.

- Observe if any logical glitches occur. Maybe you find that some economic indicator isn't available until a certain time of day and your code ran before that, causing a null, such things need handling (like coding the strategy to wait for data or carry forward last known values).
- Check if the volatility targeting is overshooting or undershooting because of estimation error. You might refine the volatility lookback period.
- Evaluate transaction costs: although ETFs are cheap to trade (zero commissions at many brokers), leveraged ETFs do have expense ratios and decay, monitor if holding them for weeks is eroding performance as expected or more.
- Make sure the strategy's **signals align with intuition**: e.g., in a paper test, if a recession started and the model went risk-off, did it actually do so? If it didn't, find out why (maybe an indicator didn't trigger as expected).
- **Human Override Plan:** Decide in advance what conditions would make you override the model. For instance, if there's a once-in-a-century event (like March 2020 pandemic crash) and the model is slow

to react due to unprecedeted data, the human manager might step in. Or if the LLM says something odd (“allocate 100% to cash for next year”), you need a process to verify or veto that.

Use this phase to also ensure the **LLM prompt process** is working, if you choose to use it in practice. For instance, if you plan to feed the strategy doc and data to GPT-4 on a schedule, test that. Check that outputs are consistent and that the cost (API calls) is justified. It might turn out simpler to just codify the decisions without the LLM, but the LLM could still be used for *explanations*. One interesting hybrid approach: run the algorithm code for actual trades, but ask the LLM to provide a plain-English explanation of the moves to include in reports. This leverages the LLM for what it’s great at (language and explanation) and code for what it’s better at (precise calculations).

6.3 Phase 3: Live Rollout

Once confident, gradually move to live trading:

- **Start Small:** Use a portion of intended capital initially. This is like a soft launch. Even if paper was fine, real markets can behave differently, and execution can have slight differences (slippage, etc.). Ensure the orders are executing as expected in real conditions.
- **Monitor Performance and Risk Daily:** At least at the start, watch the portfolio like a hawk. Confirm

that allocations match what the strategy says. Reconcile any discrepancies immediately. Keep a log of any manual interventions.

- **Real-Time Data Checks:** Make sure economic releases are being captured. For example, when CPI is released at 8:30am, do you have a mechanism to feed that number into the strategy promptly? You might have to manually input it initially or use an API like Bloomberg (expensive) or some free calendars that update (there are some free JSON feeds for econ data but reliability varies).
- **LLM Monitoring (if used):** If the LLM is part of live decision-making, consider running it in a “shadow mode” initially – i.e., get its decisions but don’t act on them until you see they consistently match your expectations. This ensures it’s parsing everything correctly. Also be mindful of any **drift** in the LLM (if it updates or changes behavior slightly over time – hopefully the deterministic nature of our prompt reduces that).
- **Communication and Reporting:** If this were a fund or a strategy for external investors, set up a monthly or quarterly report. This is where you’d use those visuals – showing how Hedge Engine is performing, explaining any deviations from benchmarks, etc. The LLM could help generate these reports by summarizing what happened (“The strategy shifted to a defensive stance in July due to weakening job growth, which helped avoid losses as the S&P fell

5%...").

6.4 Data Sources and Tools – Summary List

Finally, here's a concise list of recommended data sources and tools for ongoing use:

- **Macroeconomic Data:**

- FRED (St. Louis Fed) – API and website for thousands of series [reddit.com](https://www.reddit.com).
- BLS – website/API for labor and inflation data.
- BEA – for GDP and personal income/spending data.
- Federal Reserve websites – for Fed rates, meeting calendars, meeting minutes.
- IMF/World Bank/OECD – if expanding globally, these have international macro data (not real-time but useful for reference).

- **Market & ETF Data:**

- Yahoo Finance API or yfinance – free daily data for ETFs (prices and volumes).

- Exchanges (NASDAQ, NYSE sites) – for official closing prices.
- ETF provider websites (iShares, Vanguard, State Street): They often provide daily updated info on NAV, sometimes holdings (iShares provides downloadable holdings daily which could be useful for understanding what's inside sector ETFs, for instance).
- Financial data APIs like Alpha Vantage, IEX Cloud (for more frequent updates or extended historical).

- **Event Probabilities:**

- CME's FedWatch tool (for Fed hike probabilities derived from futures).
- Options markets (like using options implied volatility or skew as signals for event risk).
- Prediction markets (e.g., PredictIt for certain political events) – not mainstream but could provide insights into elections or policy outcomes.

- **Trading & Execution:**

- Interactive Brokers API – robust for executing trades across many assets.

- Alpaca API – a popular commission-free API broker for equities (ETFs).
 - QuantConnect or other platforms if you want to outsource some infrastructure (QuantConnect allows live trading on IB or others via their cloud with algorithms coded in Python/C#).
- **LLM/AI Tools:**
- OpenAI API (GPT-4) or similar for any AI-driven analysis or text generation (like writing commentary, or even for parsing news articles to see if any macro events need attention).
 - Python libraries for machine learning if incorporating some predictive models (scikit-learn, TensorFlow, etc., though that's beyond our rule-based approach currently).

By leveraging these resources, a retail investor can assemble an ecosystem that rivals a small hedge fund's setup: data streams, analytics, and automated execution.

6.5 Ongoing Maintenance and Evolution

The rollout doesn't end at go-live. Ongoing:

- Continuously gather performance data and compare to expectations. For example, maintain a rolling record of “predicted” vs actual outcomes for event responses (did our Fed positioning pay off or backfire? do we need to tweak how much we hedge?).
- Update the strategy if needed. The beauty of having it in a meta-prompt style is that we can edit the rules in this guide and effectively re-train the LLM or adjust the code. If a new indicator becomes important (say, after 2020, money supply or fiscal stimulus tracking became crucial), we can integrate that.
- Ensure the strategy stays within risk parameters as account values change (if capital grows, are we scaling up positions proportionally? Are there any new risks, like position sizes exceeding what the market can absorb? Unlikely with ETFs unless managing very large sums).
- Keep an eye on **regulation**: If any rules change regarding using leverage or particular ETFs (for instance, some jurisdictions have restrictions on inverse/leveraged ETFs for retail traders), comply accordingly. Always read prospectus of any ETF used – providers sometimes change how they achieve results or even close funds.
- Document everything. If this is for personal use, journaling trades and reasoning is still valuable. If

it's for clients or as a product, documentation and transparency build trust.

Conclusion:

We conclude *Hedge Engine* with a reiteration of its purpose: to bridge human-friendly investment strategy guidance with machine-executable precision. We covered how to structure information in a way that mirrors professional finance publications: clear, concise, and rich with data visualization, while also creating a blueprint an AI can directly follow. We delved into macro strategy best practices: rotating assets with the economic cycle, timing around events with probabilistic thinking, and overlaying volatility management and leverage carefully. By addressing everything from format to risk disclosures to implementation steps, this guide aimed to be comprehensive.

Ultimately, the Hedge Engine strategy empowers a sophisticated retail investor to navigate markets using a disciplined, macro-driven approach similar to those employed by hedge funds and institutional players. With prudent execution and respect for the risks involved, it can act as a “hedge fund in an ETF wrapper,” potentially enhancing portfolio resilience and returns in an ever-changing macroeconomic landscape. Use this guide as both a manual and a meta-prompt – whether you are executing the trades yourself or delegating to an AI, the principles and rules herein provide a strong foundation for macro investing success.

Addendum I: Investment Parameters and Limits

This section defines the strategy's core quantitative inputs and risk limits. These parameters are *definitive and non-negotiable*, guiding the LLM-driven model in execution and providing clear constraints for risk management:

Parameter	Value	Constraint / Reference
Target Annual Volatility	10%	Primary risk budget for the portfolio. The strategy dynamically adjusts exposure to maintain ~10% annualized volatility syntaxdata.comkundan-reads.readthedocs.io . This target level of risk is commonly used in volatility-targeting portfolios as a balance between growth and drawdown control man.com .
Volatility Measurement Window	20 Trading Days (~1 month)	Lookback period for realized volatility calculation. Using ~20 trading days of returns (approximately one calendar month) to estimate volatility provides a responsive short-term risk measure fidelity.com . This helps the model react to recent market turbulence while smoothing out daily noise.

Max Scaling Factor (Leverage)	2.0×	Hard cap on portfolio leverage to prevent over-leveraging risk. The strategy will not exceed 2x gross exposure even when volatility is very low. (Notably, U.S. regulations often impose ~2x leverage limits on funds employing volatility targeting syntaxdata.com .) This ensures the portfolio cannot take on more than double its net asset value in exposure.
Minimum Exposure Floor	0.3× (30% of full exposure)	Ensures the portfolio never fully de-risks, even under extreme volatility. A 30% minimum exposure is maintained to avoid going 100% to cash, mitigating gap risk and allowing participation in any sharp rebound ssga.com . In practice, similar risk-managed strategies set a minimum floor (often 10–25% exposure) to balance protection with opportunity cost ssga.com .
Base Rebalancing Frequency	Monthly	Standard schedule for macro rotation adjustments. The portfolio allocations are reviewed and rebalanced on a monthly cycle, which is typical for strategic sector rotation strategies pocketoption.com . (This does not preclude urgent intra-month adjustments if risk triggers are hit; it simply defines

		the routine interval for regular rebalancing.)
Emergency Drawdown Trigger	-10% (portfolio peak-to-trough)	A hard risk-off trigger. If the portfolio drops 10% from its high-water mark, an immediate reduction in risk is executed regardless of the regular schedule. This rule is designed to cap losses and preserve capital by cutting exposure during severe downturns. (For example, some traders halve or cease risk-taking after drawdowns in the ~10–15% range to prevent larger losses tradeciety.com .)
Max Single Sector Weight	25% (excl. cash/treasuries)	Diversification limit for portfolio concentration. No single equity sector (other than safe assets like cash or U.S. Treasuries) can exceed 25% of the portfolio's value. Concentrating more than a quarter of assets in one sector is generally avoided in diversified portfolios to reduce idiosyncratic risk dividendtitan.com . This cap ensures the strategy maintains broad exposure across sectors.

Addendum II: Definitive ETF Universe and Tickers

This section enumerates the approved universe of Exchange-Traded Funds (ETFs) that the strategy can trade. Each asset category is tied to a specific macroeconomic role or cycle phase, ensuring the LLM executes trades on liquid instruments that accurately represent the intended exposure. The list is **unambiguous** – these tickers define exactly what assets the model may rotate into, aligning with the strategy's market regime framework:

Asset Category	Target Exposure Role	Example Tickers (Illustrative)
Broad Equity Market	Baseline Core Exposure (Market Beta)	SPY, VTI – U.S. large-cap S&P 500 ETF (SPY); total U.S. stock market ETF (VTI)
Cyclical Sectors <i>Tech, Discretionary, Industrials, Financials</i>	Overweight in Early-Cycle expansions (risk-on growth)	XLK, XLY, XLI, XLF – Sector SPDRs for Technology, Consumer Discretionary, Industrials, Financials
Defensive Sectors <i>Staples, Healthcare, Utilities</i>	Overweight in Late-Cycle or Recession (risk-off defensives)	XLP, XLV, XLU – Sector SPDRs for Consumer Staples, Health Care, Utilities
Long-Term Bonds	Defensive Recession Hedge	TLT, IEF – iShares 20+ Year Treasury Bond ETF (TLT);

		<i>7–10 Year Treasury ETF (IEF)</i>
Inflation/Commodity Hedge	Late Cycle / Inflation Regime Hedge	GLD, DBC – <i>SPDR Gold Trust (GLD); Invesco DB Commodities Index ETF (DBC)</i>
Cash Proxy / Short-Term Bonds	Risk-Off / Capital Preservation	SHY, BIL – <i>iShares 1–3 Year Treasury Bond ETF (SHY); SPDR Bloomberg 1-3 Month T-Bill ETF (BIL)</i>
Approved Leveraged Instruments	For Volatility Scaling > 1.0× (when leverage is applied)	SSO, TBT – <i>ProShares Ultra S&P 500 (SSO, 2× leveraged S&P 500); ProShares UltraShort 20+ Year Treasury (TBT, –2× inverse Treasuries)</i>

Note: The asset categories and their target uses reflect a **business cycle rotation** approach. In early economic expansions, **cyclical sectors** (e.g. technology, consumer discretionary, industrials, financials) tend to outperform and thus are overweighted [beaconinvesting.com](#). Conversely, in downturns or recessions, **defensive sectors** (staples, utilities, healthcare) hold up better and receive higher allocation [beaconinvesting.com](#). Broad market ETFs like SPY/VTI provide core equity exposure, while Treasury bond ETFs (TLT, IEF) and cash proxies (SHY, BIL) serve to mitigate risk in risk-off periods. Additionally, modest use of **leveraged ETFs** (such as SSO and TBT) is permitted *only* to

achieve the volatility target when needed – for instance, adding 2× S&P 500 exposure via SSO in low-volatility environments, or using TBT to hedge interest rate risk – and is constrained by the 2.0× leverage cap noted in Addendum I.

Addendum III: Backtest Performance Summary (Placeholder)

(To be updated once strategy backtesting is completed) – The table below is a template for comparing the **Hedge Engine Strategy**'s performance against a traditional 100% equity benchmark (S&P 500 Index) and a classic 60/40 stock-bond portfolio. The goal is to demonstrate that the strategy achieves **comparable returns with reduced drawdowns and volatility**, resulting in a smoother performance profile than the benchmarks stanlib.com. Key metrics like risk-adjusted returns (Sharpe ratio) and Calmar ratio are expected to be superior for the Hedge Engine, reflecting its ability to limit losses. **All figures will be validated via historical backtesting:**

Metric	Hedge Engine Strategy (Target Outcome)	S&P 500 (Benc hmark)	60/40 Portfolio (Bench mark)
Compound Annual Growth Rate (CAGR)	Competitive with Benchmark (≈ match S&P)	Placeholder r	Placeholder r
Max Drawdown Worst Peak-to-Trough Loss	Significantly Lower (shallower declines)	Placeholder r (e.g. -50%)	Placeholder r
Sharpe Ratio Risk-Adjusted Return	Higher (better risk-adjusted performance)	Placeholder r	Placeholder r

Calmar Ratio	Significantly Higher	<i>Placeholder</i>	<i>Placeholder</i>
CAGR / Max Drawdown	(smoother returns)	r	

Explanation: The Hedge Engine Strategy is designed to **reduce volatility and drawdowns** relative to a full-equity portfolio, while maintaining similar average returns over the long runstanlib.com. A successful backtest would likely show the strategy with a much smaller maximum drawdown than the S&P 500 (for context, the S&P 500's peak-to-trough drawdown in the 2008 crisis was roughly -50%cdn.janushenderson.com, whereas a 60/40 portfolio fell by about -35% in the same period). By dynamically scaling risk, the strategy aims to avoid such large losses – improving the Sharpe ratio (higher reward per unit of volatility) and boosting the Calmar ratio (higher return per unit of drawdown). These outcomes would validate the strategy's premise of **smoother returns and improved capital preservation**stanlib.com, which are crucial for compounding wealth and maintaining investor confidence during market stress. All performance figures will be updated with actual backtested data, replacing the placeholders above, to provide verifiable evidence of the strategy's efficacy once analysis is completed.

Addendum IV — Auditable Functionality

Purpose

This addendum defines the auditable functionality required for Hedge Engine. It describes the mandatory logging, provenance, governance, monitoring, testing, and report formats that make every automated decision traceable, reproducible, and defensible for internal review and external audit.

Scope

Applies to:

- All LLM prompts, prompt versions, and system prompts used to generate trade signals.
- All deterministic code (validation, decay simulation, position sizing, execution logic).
- All human approvals, overrides, and escalations.
- All execution messages sent to brokers, order managers, or downstream systems.

Principles (non negotiable)

1. Every decision must be reconstructible. Given the record you must be able to replay the inputs and arrive at the same numeric signal and execution recommendation.
2. Evidence must accompany probabilistic statements. Any P success or expected delta must cite the underlying

sources the LLM used.

3. Immutable provenance. Prompt versions, model id, and evidence must be recorded in a write once log.
4. Human oversight when confidence is low or the decision crosses policy thresholds. Low confidence triggers an approval gate.
5. Retention and accessibility. Audit records must be retained for regulatory period (default 7 years) and be searchable.

Required Components

1) Decision Record (atomic unit)

Every decision or signal must produce one Decision Record. This is the canonical artifact auditors read.

Decision Record JSON Schema (required fields):

```
{  
  "decision_id": "uuid",  
  "timestamp_utc": "ISO8601",  
  "model_version": "string",  
  "prompt_hash": "sha256",  
  "prompt_text": "string",  
  "inputs": { "market_data_snapshot": {...}, ... },  
  "macro_inputs": {...}, "event_calendar": {...}  
},  
  "evidence": [
```

```
    { "source_id": "string",
      "type": "CRS|docket|news|options_skew|ptr",
      "filecite": "file_id|L10-L20", "excerpt": "string"
    }
  ],
  "llm_output": {
    "structured_schema_version": "1.0",
    "p_success": 0.0,
    "p_confidence": 0.0,
    "horizon_days": 0,
    "expected_delta": null
  },
  {"fav": 0.0, "neutral": 0.0, "unfav": 0.0},
  "suggested_instrument": {
    "type": "LETF|ETF|Futures|Options|Swap",
    "tactic": "core|tactical"
  },
  "quant_checks": {
    "ev_gross": 0.0,
    "letf_decay": 0.0,
    "ev_net": 0.0,
    "viability_pass": true
  },
  "human_review": {
    "required": true,
    "reviewer_id": "string",
    "approval": "approved|rejected|escalated",
    "notes": "string"
  },
  "execution_plan": {
```

```
    "orders":  
    [ {"instrument": "SPY", "side": "BUY", "qty": 123, "type": "LIMIT|MKT", "timestamp_planned": "iso"}],  
      "sor_policy": "percent_of_adv=10"  
    },  
    "audit_hash": "sha256-of-entire-record",  
    "signature": "system|user-sig"  
}
```

Notes:

- `prompt_hash` and `model_version` must reference the exact LLM model and prompt text used.
- `evidence` must include at least two distinct source hits for any high-conviction trade (CRS doc plus options skew, or docket plus news, etc.).
- `quant_checks` shows the deterministic validation that gates execution (EV calculation, decay simulation). Execution is forbidden unless `viability_pass` is true or a documented human override occurs.

2) Prompt and Model Versioning

- Every prompt is versioned in source control with a human readable changelog.
- Prompt versions are immutable once used for live decisions; any edits must create a new version.

- Model id and provider are recorded (example `gpt-4o-2025-11-15-v2`). Any change in model provider triggers full revalidation of prompts.

3) Evidence Requirements

- For any probabilistic claim the LLM makes, evidence array must contain: source id, paragraph or lines referenced, citation timestamp, and excerpt.
- If LLM cites internal proprietary data, the dataset id and hash must be included so auditors can verify integrity.

4) Deterministic Gates and Tests

- Expected Value Gate and Decay Simulation must be implemented in deterministic code that is unit tested. Tests must cover edge cases: low liquidity, gap scenario, worst case decay.
- Viability tests must be deterministic and reproducible by auditors using the recorded `inputs` from the Decision Record.

5) Execution Audit Trail

- Every order executed must link to the Decision Record id. The execution log must contain order id, time, fills, venue, slippage, and final P&L attribution.
- If an execution deviates (manual split, partial fills), the operator must annotate the Decision Record with reason

and reconciliation.

6) Human in the Loop

- Human approval is mandatory when any of the following occur:
 - `p_confidence < 0.7` and `ev_net > safety_margin`
 - Suggested use of LETF for `horizon_days > T_max` (the decay horizon)
 - Total gross leverage post trade would exceed portfolio cap
- Approval is captured in `human_review` block with timestamp, reviewer id, and signed note.

7) Monitoring and Drift Detection

- Telemetry must track: mean pass rate of viability gate, hallucination flags per model, distribution of `p_success` vs realized outcomes, and evidence completeness rate.
- Alerts for anomalies:
 - Hallucination rate > threshold (e.g., 5% of outputs missing required evidence)

- Viability pass rate falling dramatically after model updates
- Model version change without revalidation
- Drift detection: compare rolling calibration of p_success to realized event frequency. If model overstates probability by more than X points over 90 days, freeze related prompts and require retrain.

Audit Reports and Artifacts

A) Daily Operational Report (automated)

- Number of decisions made
- Decisions with `human_review.required=true` and their status
- Number of orders executed, total notional, avg slippage
- Hallucination rate, worst offenders (prompt hashes), top evidence failures
- Alerts triggered and remediation actions

B) Monthly Compliance Pack (for auditors)

- Index of Decision Records for the month (`decision_id` list)
- Copies of all prompt versions and model ids in use

- Code hash of deterministic gating code (EV Gate, Decay Simulation)
- Sample replay: reproduction instructions to reconstruct 3 random decisions end to end
- Incident log and resolution record

C) Ad hoc Forensic Replays

- The system must be able to replay any Decision Record end to end including: loading market snapshot, re-running LLM call with original prompt and model id (if provider supports), re-running deterministic gates, and re-simulating orders.

Operational Procedures

1) Onboarding a New Prompt

1. Author writes prompt and JSON output schema.
2. Unit tests and scenario tests created (including adversarial cases).
3. Validation run on historical events and synthetic edge cases.
4. Security review and privacy check.
5. Version signed and deployed to staging.

6. Pilot period with shadow decisions only.
7. Promote to live after performance criteria met.

2) Emergency Response Flow

- If a serious incident occurs (mispriced trade, model hallucination causing loss > threshold), follow this helpline:
 1. Kill switch: pause automated execution.
 2. Triage: collect Decision Records, execution logs, model id.
 3. Human review and patch.
 4. Re-run affected decisions in replay mode and produce root cause report.
 5. Restore only after remediation and revalidation; produce compliance report.

Metrics for Governance

- **Audit coverage:** percent of automated actions with a complete Decision Record. Target 100%.
- **Reproducibility time:** time required to fully reconstruct a Decision Record. Target < 24 hours.

- **Evidence completeness:** percent of decisions meeting the evidence minimum. Target > 99%.
- **Mean time to detect drift:** Target < 7 days.
- **Human override rate:** percent of decisions requiring human approval. Track trend.

Privacy, Retention, and Access Control

- Audit logs retained minimum 7 years, encrypted at rest.
- Access control: role based access. Compliance and auditors have read only rights. Only approved operators can mark `human_approved`.
- Sensitive PII must be redacted from logs and stored separately with stricter governance.

Example Audit Entry (human friendly summary)

- Decision id: `uuid`
- Date: 2025-12-01 14:05:00 UTC
- Event: FOMC probabilities triggered, suggested reduce Treasury duration
- LLM model: gpt-fin-2025-11-v3 prompt hash: `abcd1234`

- Evidence: Fed minutes (filecite:R.12|L12-L26), market implied prob from CME (filecite:CME.2025-12-01)
- EV net: +3.8% viability pass: true
- Human reviewer: jane.doe@example.com Approved at 14:18 UTC
- Order executed: reduce TLT 20% to IEF, fill slippage 0.03% total notional \$5M
- Audit hash: xyz987

Implementation Checklist (minimum)

- Decision Record schema implemented and enforced for all decisions
- Immutable log (append only) with cryptographic hashing of records
- Prompt versioning and model id recording in source control
- Deterministic EV Gate and Decay Simulation tests committed, with code hashes recorded
- Human in loop gates implemented and enforced by policy
- Replay functionality verified for random sample of decisions
- Daily and monthly audit reports automated

- Drift detection and alerting configured
- Incident response runbook tested

Final guidance

This addendum belongs directly after the Conclusion in every published Hedge Engine deliverable. It must be treated as living policy: any change to prompts, model versions, or gating code requires an update to this addendum and a new signed revision. If you want, I will convert this addendum into a machine-checkable compliance manifest (JSON+YAML) and a sample audit pack containing three example Decision Records ready for an auditor.

Addendum V — Custom GPT Configuration (for a Hedge Engine custom GPT)

Purpose

Turn the Hedge Engine manuscript into a reproducible, auditable *custom GPT* that acts as the strategy's knowledge base and rules engine. This config forces structured outputs, enforces evidence, and hands deterministic gates to downstream systems.

1) System message (copy/paste)

SYSTEM: You are HEDGE ENGINE GPT – an expert, auditable macro ETF strategy agent.

Behave as a disciplined quantitative analyst and compliance-first trading assistant. Output only valid JSON matching the required schema. Always include evidence citations for probabilistic claims. Never produce free-format trading orders. If confidence < 0.7, output "requires_human_review". Do not attempt execution. Log prompt_hash and model_version in every response.

2) Required structured output (JSON schema — single-response)

```
{  
    "decision_id": "string",  
    "timestamp_utc": "ISO8601",  
    "context": "{cycle_phase, event_calendar,  
    market_snapshot}" ,
```

```
"llm_signal":{  
    "p_success":0.0,  
    "p_confidence":0.0,  
    "horizon_days":0,  
  
    "expected_delta":{"fav":0.0,"neutral":0.0,"unfav":0.0},  
    "trade_strength":-1.0,  
  
    "suggested_instrument":{"ticker":"string","type":"ETF|LETF|FUT|OPT","tactic":"core|tactical"}  
},  
  
"evidence":[{"source":"string","id":"string","excerpt":"string","filecite":"string"}],  
    "explainability":["short bullet rationales"],  
  
    "flags":{"requires_human_review":false,"low_evidence":false,"out_of_universe":false},  
    "prompt_hash":"string",  
    "model_version":"string"  
}
```

Enforcement: any response that fails schema should be rejected by the runtime. Do not accept plain text.

3) Few-shot / Examples (short)

Include 2–3 structured examples inside the custom GPT: one early-cycle tilt, one Fed-event hedge, one low-confidence output

that triggers human review. Keep these in the prompt repository as `examples/v1.json`.

4) Prompt patterns for operations (safe templates)

Signal generation

USER: SIGNAL_GENERATION

Context: {`market_snapshot`, `macro_indicators`,
`options_skew`, `ptr_swarm`}

Task: Return JSON per schema. Cite at least 2 evidence items for `p_success>=0.7`.

-

Evidence augmentation

USER: EVIDENCE_AUGMENT

Provide 3 source ids that support the LLM's top claim and a single-sentence excerpt for each.

-

5) Tooling & connectors

- Required connectors: market price API, options chain API, CRS/legislative documents, PTR filings, news stream, internal model store.
- Provide a deterministic wrapper that fetches data snapshots and injects `market_snapshot` into the prompt. Include timestamps, exchange ids, and liquidity metrics.

6) Safety / compliance gates baked in

- If `p_confidence < 0.7` → set `flags.requires_human_review=true`.
- If suggested instrument not in Definitive ETF Universe → set `flags.out_of_universe=true`.
- If `horizon_days > T_max` for LETF but `suggested_instrument.type == LETF` → set `flags.low_evidence=true` and require human approval.

7) Testing and validation

- Unit tests: schema validation; hallucination checks (evidence completeness); probabilistic calibration tests (calibrate `P_success` on holdout events).
- Scenario tests: replay 50 historical events and assert `viability_pass` rates and evidence completeness.

8) Versioning & provenance

- Save: `prompt_text`, `prompt_hash` (SHA256), `model_version`, `date`, and `example outputs` every deployment. Store in an immutable repo. Require signed change log for any prompt edits.

Addendum VI — Operator Instructions to Make Hedge Engine *Agentic*

Purpose

Guide the operator to safely instantiate Hedge Engine as an agentic system — autonomous in planning and execution but governed by auditable controls, human gates, and a kill switch. This is not “let it trade by itself forever.” It’s safe autonomy: planning, proposal, deterministic checks, pre-commit human approvals for risky ops.

1) Agent architecture (high level)

- **Planner:** Breaks the high-level objective into tasks (e.g., adjust sector tilt, set hedges, manage vol). Uses LLM for strategy reasoning and scenario simulation.
- **Controller / Executor:** Deterministic engine that builds orders and applies execution policy (SOR, participation limits). It **never** places orders without the **Execution Gate** passing.
- **Validation Layer:** EV Gate + Decay Simulation + PositionSizing module (deterministic code). Failsafe before execution.
- **Audit Ledger:** Append-only store of Decision Records and Execution Logs.
- **Human Ops Console:** Displays proposals requiring human sign-off and emergency controls.

- **Monitoring and Observability:** Real-time dashboards for exposures, leverage, drift, and hallucination metrics.
-

2) Agentic loop (safe pseudocode)

```
while trading_hours:  
    snapshot = fetch_market_and_macro_snapshot()  
    plan = planner.generate_plan(snapshot) #  
    returns list of actions  
    for action in plan:  
        decision_record =  
    llm_signal(action.context) # structured JSON  
        record(decision_record)  
  
        # Deterministic validation  
        checks =  
        run_viability_checks(decision_record)  
        if not checks.viability_pass or  
decision_record.flags.requires_human_review:  
            route_to_human(decision_record)  
            continue  
  
            # Pre-exec orchestra  
            execution_plan =  
controller.build_execution(decision_record)  
            if execution_plan.exceeds_limits():  
                escalate_and_log(decision_record,  
execution_plan)  
                continue
```

```

# Optional: fast-path for small, safe
adjustments
    if execution_plan.is_low_risk():
        controller.execute(execution_plan)
        log_execution(execution_plan,
decision_record)
    else:
        # require explicit operator approval
        mark_for_approval(decision_record,
execution_plan)

```

3) Operator approvals & policy thresholds

Operators configure **policy profiles** that define when agent may auto-execute vs require manual signoff. Example conservative policy:

- Auto-exec allowed if: notional per trade < \$50k, gross portfolio leverage change < 0.05, p_confidence \geq 0.85, ev_net > 1.5%
- Manual approval if: any LETF > 0.5% of NAV, horizon_days > 7 and uses LETF, or p_confidence between 0.7 and 0.85.

Approval UI must show: Decision Record, evidence excerpts, deterministic EV math, decay simulation, proposed orders, estimated slippage, and counterparty constraints.

4) Kill switch, sandboxing, and canaries

- **Global Kill Switch:** Hard stop that turns agentic actions to shadow mode (no live order). Must be hardware or cloud-managed with two-person custody for reactivation.
 - **Sandbox Mode:** Agent can run in a market-simulated environment using real-time prices and a simulated brokerage for stress testing. Use this for model updates.
 - **Canary Execution:** When promoting a new agent behavior, allow a canary window: small notional execution with expanded monitoring and rapid rollback rules.
-

5) Execution policy: safety-first defaults

- **Max per-trade notional** = $\min(\text{configured_limit}, X\% \text{ NAV})$
- **Participation rate** = max 1% to 10% of ADV configurable by instrument liquidity
- **Order types** = prefer limit orders, TWAP/VWAP for big blocks
- **Slippage guard** = abort if estimated slippage > allowed bps threshold

- **Post-fill reconciliation** = link fills to Decision Record and compute realized ev attribution
-

6) Monitoring, drift detection, calibration

- Track rolling calibration: compare predicted p_success vs realized frequency; trigger retrain or freeze if miscalibration > threshold over 30–90 days.
 - Monitor hallucination: fraction of outputs lacking required evidence. Auto-freeze prompts with evidence failure rate > 2–5%.
 - Daily health check: viability_pass rate, human review queue size, avg time to approve, total exposures vs policy.
-

7) Incident playbook (operator steps)

1. **Detect**: alert triggers (slippage, hallucination, execution anomaly).
2. **Kill**: immediately flip to shadow mode.
3. **Collect**: gather Decision Records, execution logs, model_version, prompt_hash.

4. **Triage**: quick root cause (data error, model change, prompt bug, broker outage).
 5. **Contain**: revert positions if needed (preapproved unwind script).
 6. **Remediate**: patch prompt/code and run sandbox replay tests.
 7. **Report**: produce forensic report for compliance and stakeholders.
-

8) Governance & human oversight rules (must-haves)

- Require two-person review for system changes that modify prompts, model versions, or deterministic gating code.
 - Keep a signed change log and operational playbook.
 - Enforce quarterly red-team: adversarial scenarios to find failure modes.
 - Maintain a compliance interface exporting monthly audit packs.
-

9) Practical deployment checklist

- Deploy the Planner & LLM in staging; run 60-day shadow mode with replay of historical month-by-month.
 - Validate Decision Record reproducibility for a random sample of 100 decisions.
 - Configure operator policies and test UI approval flows.
 - Test the kill switch and canary execution 3 times with simulated fills.
 - Launch with small AUM; expand after 30 days of stable metrics.
-

Final reality check

Making Hedge Engine agentic is feasible and powerful, but **never** run fully autonomous on real capital without hardened governance, exhaustive replay testing, and real human-in-the-loop approval for high-risk actions. Agentic systems accelerate operations — they also accelerate mistakes. Keep auditable records, strict thresholds, and an immutable trail. That is the only way to be both agentic and responsible.

